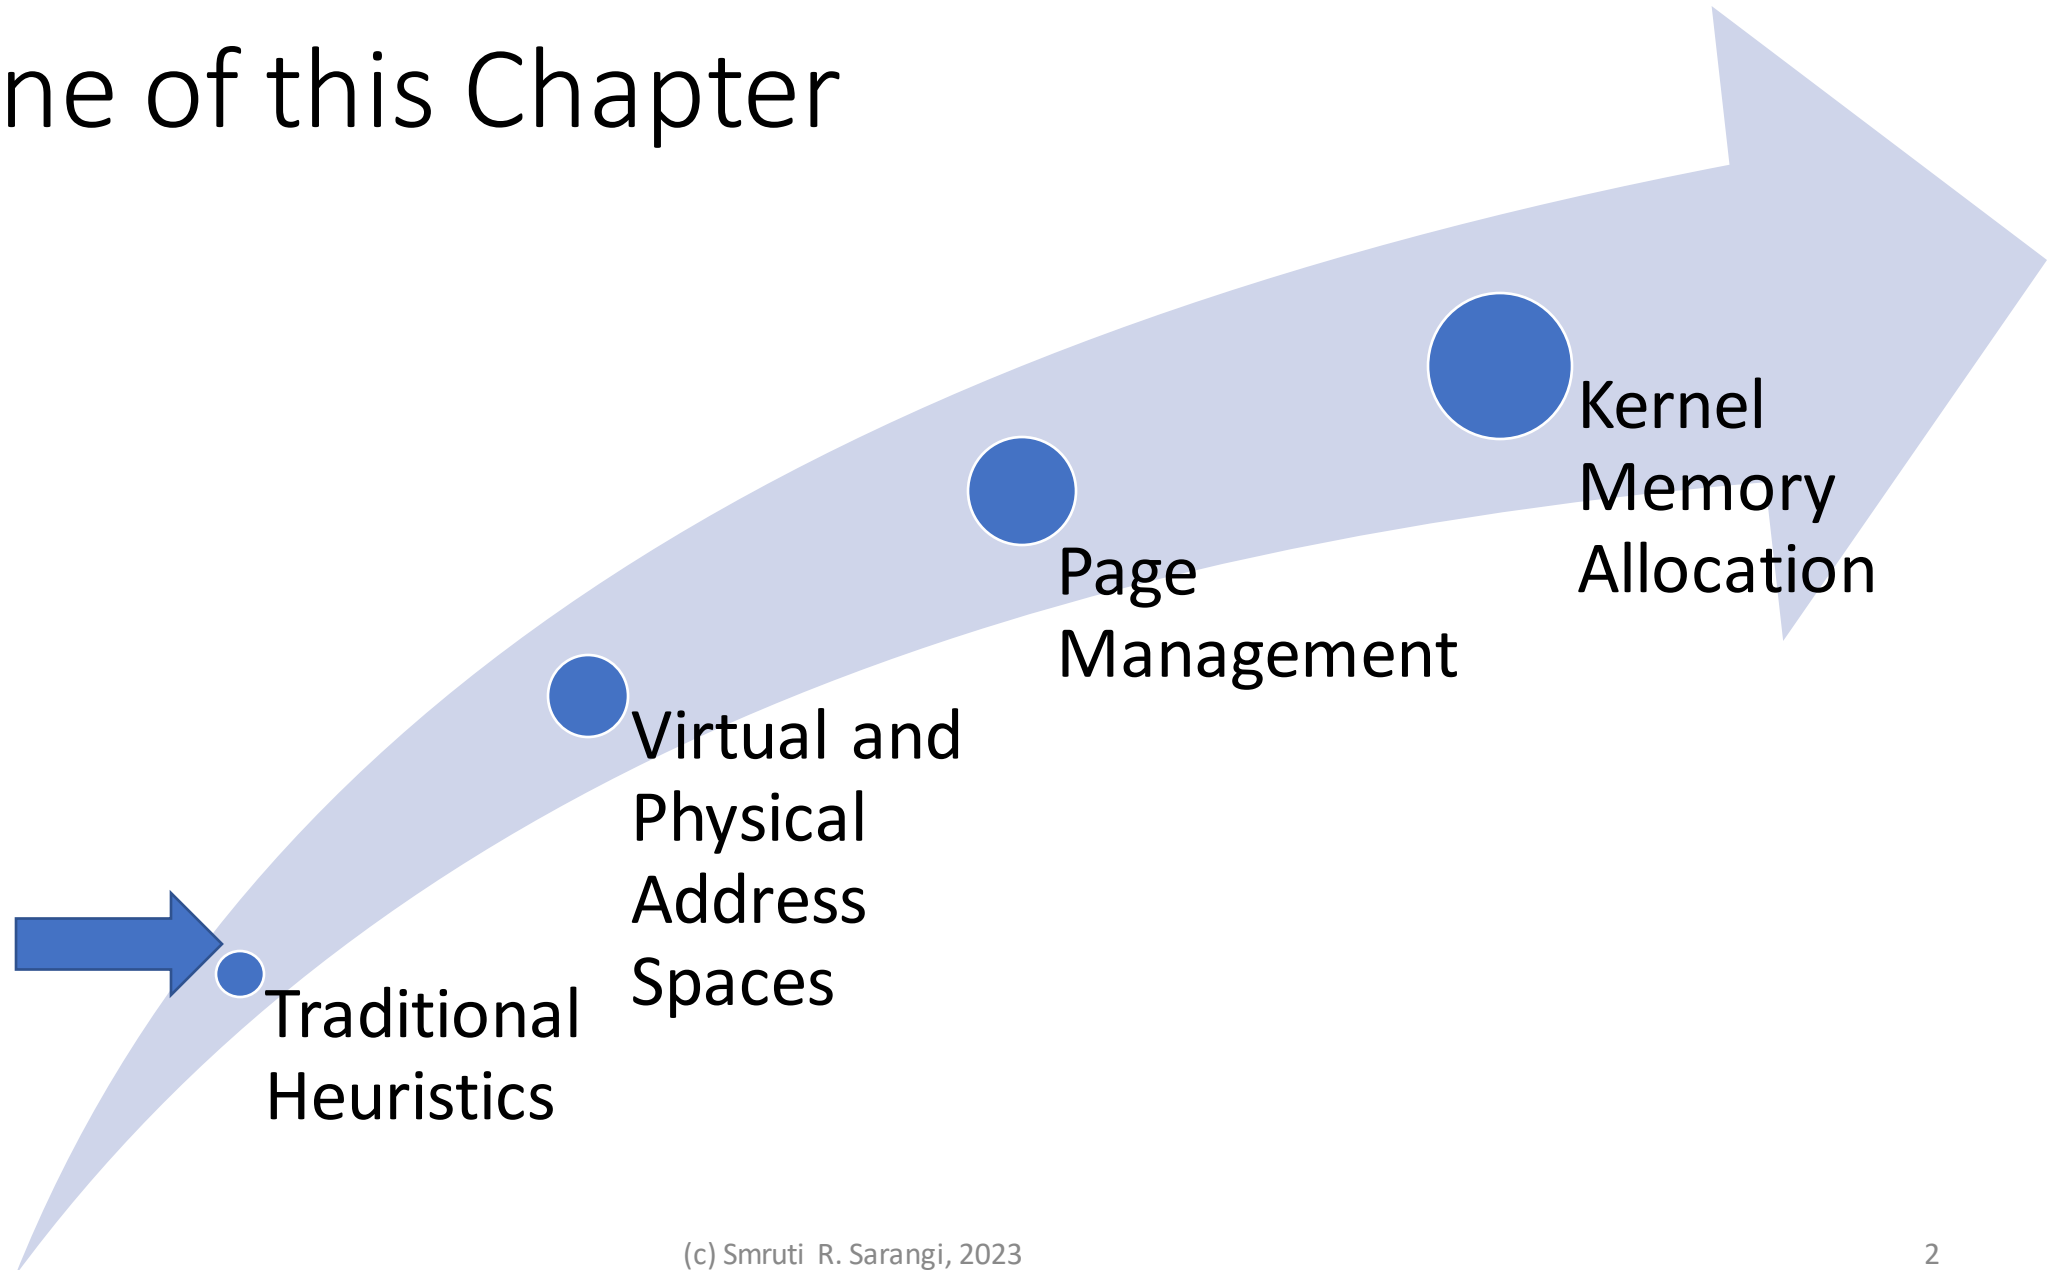


Chapter 6: Memory Systems

Smruti R Sarangi
IIT Delhi

Outline of this Chapter

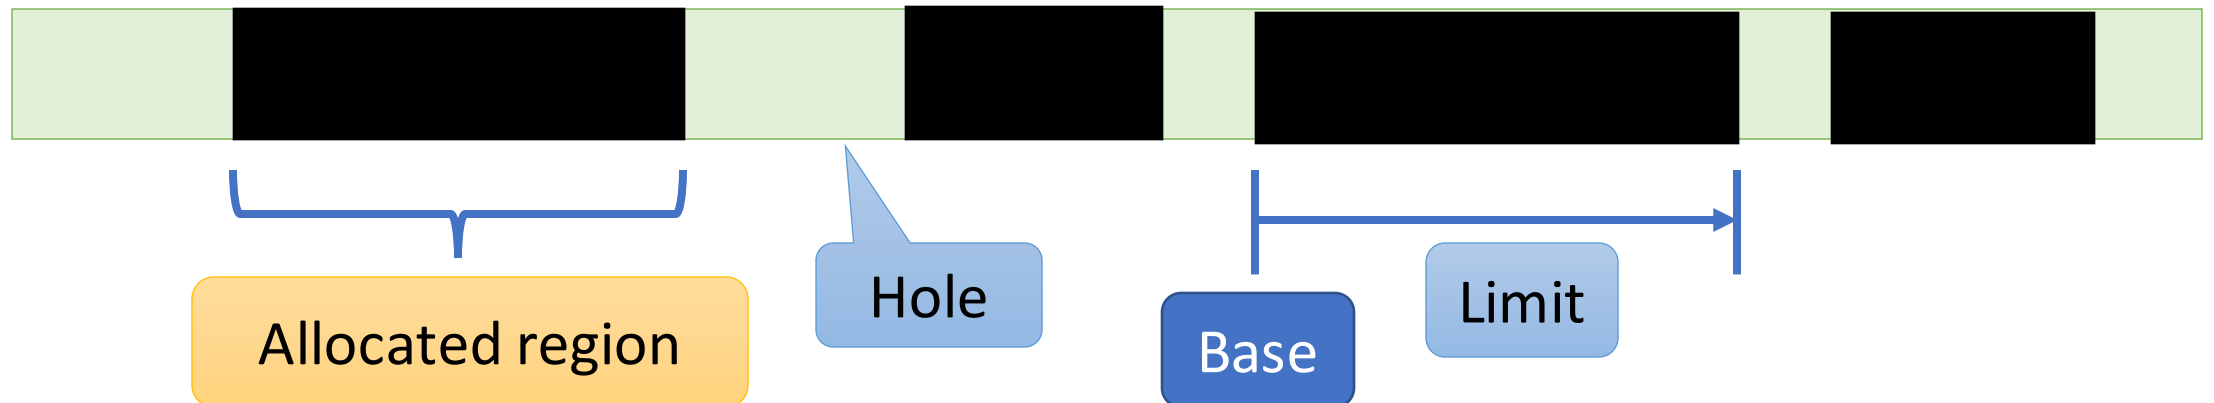


Heuristics for memory allocation: no virtual memory

Simple Memory Allocation

- Consider the **era** that did not have virtual memory
- OR systems that don't have **virtual memory**
- OR the parts of the **kernel** that need to use physical memory

Memory



Fragmentation



Space wastage

Internal Fragmentation

- Space **wasted** within an allocated region. Let us say that we waste the last 4 KB in a region.

External Fragmentation

- **Holes** between **regions**.

Algorithms to Allocate Space

- Let's say that there is a **request** for a **memory** region **R**

Memory



Which hole do we fill?

Several Solutions

Algorithms to Allocate Space – II

| | |
|-----------|---|
| Best Fit | <ul style="list-style-type: none">• Choose the smallest hole that is just about larger than R. |
| Worst Fit | <ul style="list-style-type: none">• Choose the largest hole. |
| Next Fit | <ul style="list-style-type: none">• Start searching from the last allocation that was made towards higher addresses (with wraparounds). |
| First Fit | <ul style="list-style-type: none">• Choose the first available hole |


Heuristics for memory allocation: with virtual memory

? What about allocating free frames?

- We can only have **internal fragmentation** with virtual memory
- Use a **bitmap** or an **Emde Boas tree** to manage the list of free frames
- The main issue here is **page replacement**



If the **main memory** is **full**, which frame in **memory** should be sent to the swap space?




General Stack-based Algorithms

The **notion** of the stack distance

Hypothetically


1. **Maintain** a stack of all page accesses
2. Whenever there is a page access, **locate** the **entry** in the stack.
3. **Record** the distance from the top of the stack → the **stack distance**.
4. Bring it to the **top**.



Typical Stack Distance Plot

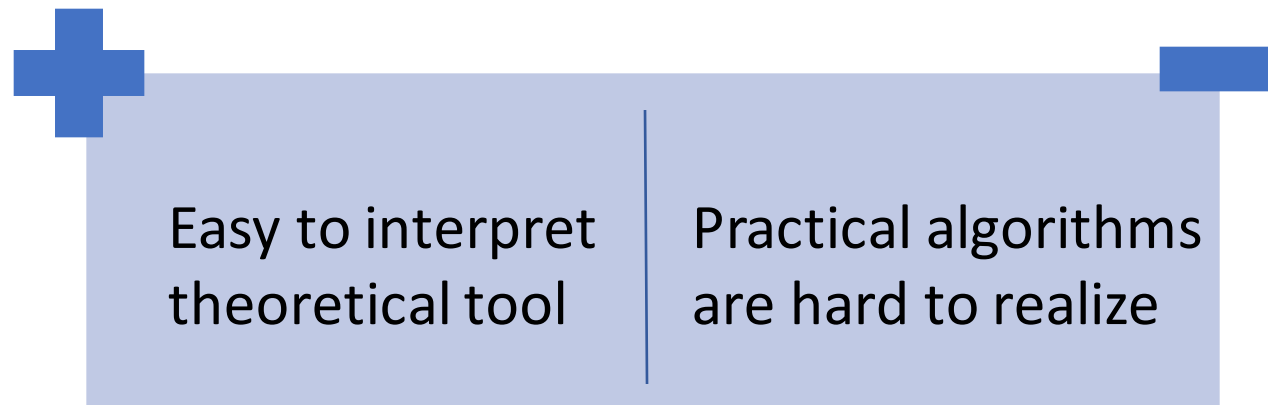
Example

A heavy tailed curve. Matches a log-normal distribution (most of the time)



Significance of the Stack Distance

- It is a **measure** of temporal locality
- **Higher** is the stack distance, **lower** is the temporal locality and vice versa
- The log-normal curve **implies** the following:
 - There are very few **accesses** with ultra-small stack **distances**
 - There is a **distinct** peak and a **long** tail



Optimal Algorithm

- Cost of a **page replacement** algorithm = Number of **page** faults

Hypothetical Optimal Algorithm

1. Order all the **pages** in increasing order of “next use” time. Assume you can **predict** the future.
2. **Replace** the page that will be accessed the **last**.



Use the same **contradiction-based** technique to prove **optimality**.

This is a stack-based algorithm. The replacement is done based on the stack distance.

The LRU (Least Recently Used) Algorithm

Impractical !!!

- Tag each page (in the memory) with the last time that it was accessed.
- Augment each entry in the page table with a timestamp.
- Choose the page whose last access time is the earliest
- Let us assume that the past is a good predictor of the future.
- These algorithms come with many theoretical guarantees




Every memory access cannot increment a counter. The overheads are prohibitive.

Such Stack-based Algorithms can be made Practical


- **Order** all the physical pages (frames) in memory in some order (maybe **FIFO** order)
- **Leverage** their protection bits. Mark them “no access”. **Access bit = 0**
- When we **access** a page with its **access** bit set to 0, there is a fault
 - **Set** it to 1
- **Periodically**, set all access bits to 0
- OR, **alternatively** record the time that it was set from 0 → 1. Reset it to 0 only if a certain **duration** has been exceeded.

WS_Clock Page Replacement Algorithm



- A **pointer** (like the **minute** hand of a clock) **points** to a physical page
- **If** its (access bit = 1) set it to 0
- Otherwise, **replace** it
- Next time, start the **pointer** from the **same** point and **wraparound** at the end

WS_Clock based Second Chance Algorithm



| <Access bit, Modified bit> | New state | Action |
|-------------------------------|-----------|---|
| <0, 0> | <0, 0> | Go ahead and replace |
| <0, 1> | <0, 0> | Schedule a write-back |
| <1, 0> | <0, 0> | Move forward |
| <1, 1> | <1, 0> | Frequently used frame; move forward. Schedule a write-back. |

FIFO and the Belady's Anomaly [1]

10 faults

Consider the FIFO page replacement algorithm

Access
sequence

| | | | | | | | | | | | |
|--------------|--------------|--------------|--------------|----|----|--------------|--------------|--------------|--------------|--------------|--------------|
| 1 | 2 | 3 | 4 | ✓1 | ✓2 | 5 | 1 | 2 | 3 | 4 | 5 |
| 1 | 2 | 3 | 4 | 4 | 4 | 5 | 1 | 2 | 3 | 4 | 5 |
| | 1 | 2 | 3 | 3 | 3 | 4 | 5 | 1 | 2 | 3 | 4 |
| | | 1 | 2 | 2 | 2 | 3 | 4 | 5 | 1 | 2 | 3 |
| | | | 1 | 1 | 1 | 2 | 3 | 4 | 5 | 1 | 2 |













4 frames

FIFO and the Belady's Anomaly

Consider the FIFO page replacement algorithm

9 faults

Access
sequence

| | | | | | | | | | | | |
|---|---|---|--|---|---|---|---|---|---|---|---|
|  |  |  |  |  |  |  |  |  |  |  |  |
| 1 | 2 | 3 | 4 | 1 | 2 | 5 | 1 | 2 | 3 | 4 | 5 |
| 1 | 2 | 3 | 4 | 1 | 2 | 5 | 5 | 5 | 3 | 4 | 4 |
| | 1 | 2 | 3 | 4 | 1 | 2 | 2 | 2 | 5 | 3 | 3 |
| | | 1 | 2 | 3 | 4 | 1 | 1 | 1 | 2 | 5 | 5 |

3 frames

Final Word on Page Replacement

- Stack-based algorithms are by and large **impractical**
- We need to create **approximations**.
- FIFO and Random **replacement** algorithms may exhibit the Belady's anomaly
 - The **ratio** between the **faults** in a large memory and small memory is unbounded \rightarrow can be as large as we want it to be [2].
- The clock-based **algorithms** approximate LRU and are known to **work** well.

Working Sets



Assume you have a recency-based replacement algorithm


- A **working set** is a set of pages that a program access over a **short duration**.



How **short** is **short**?



Keep track of the working set and ensure that it is in memory

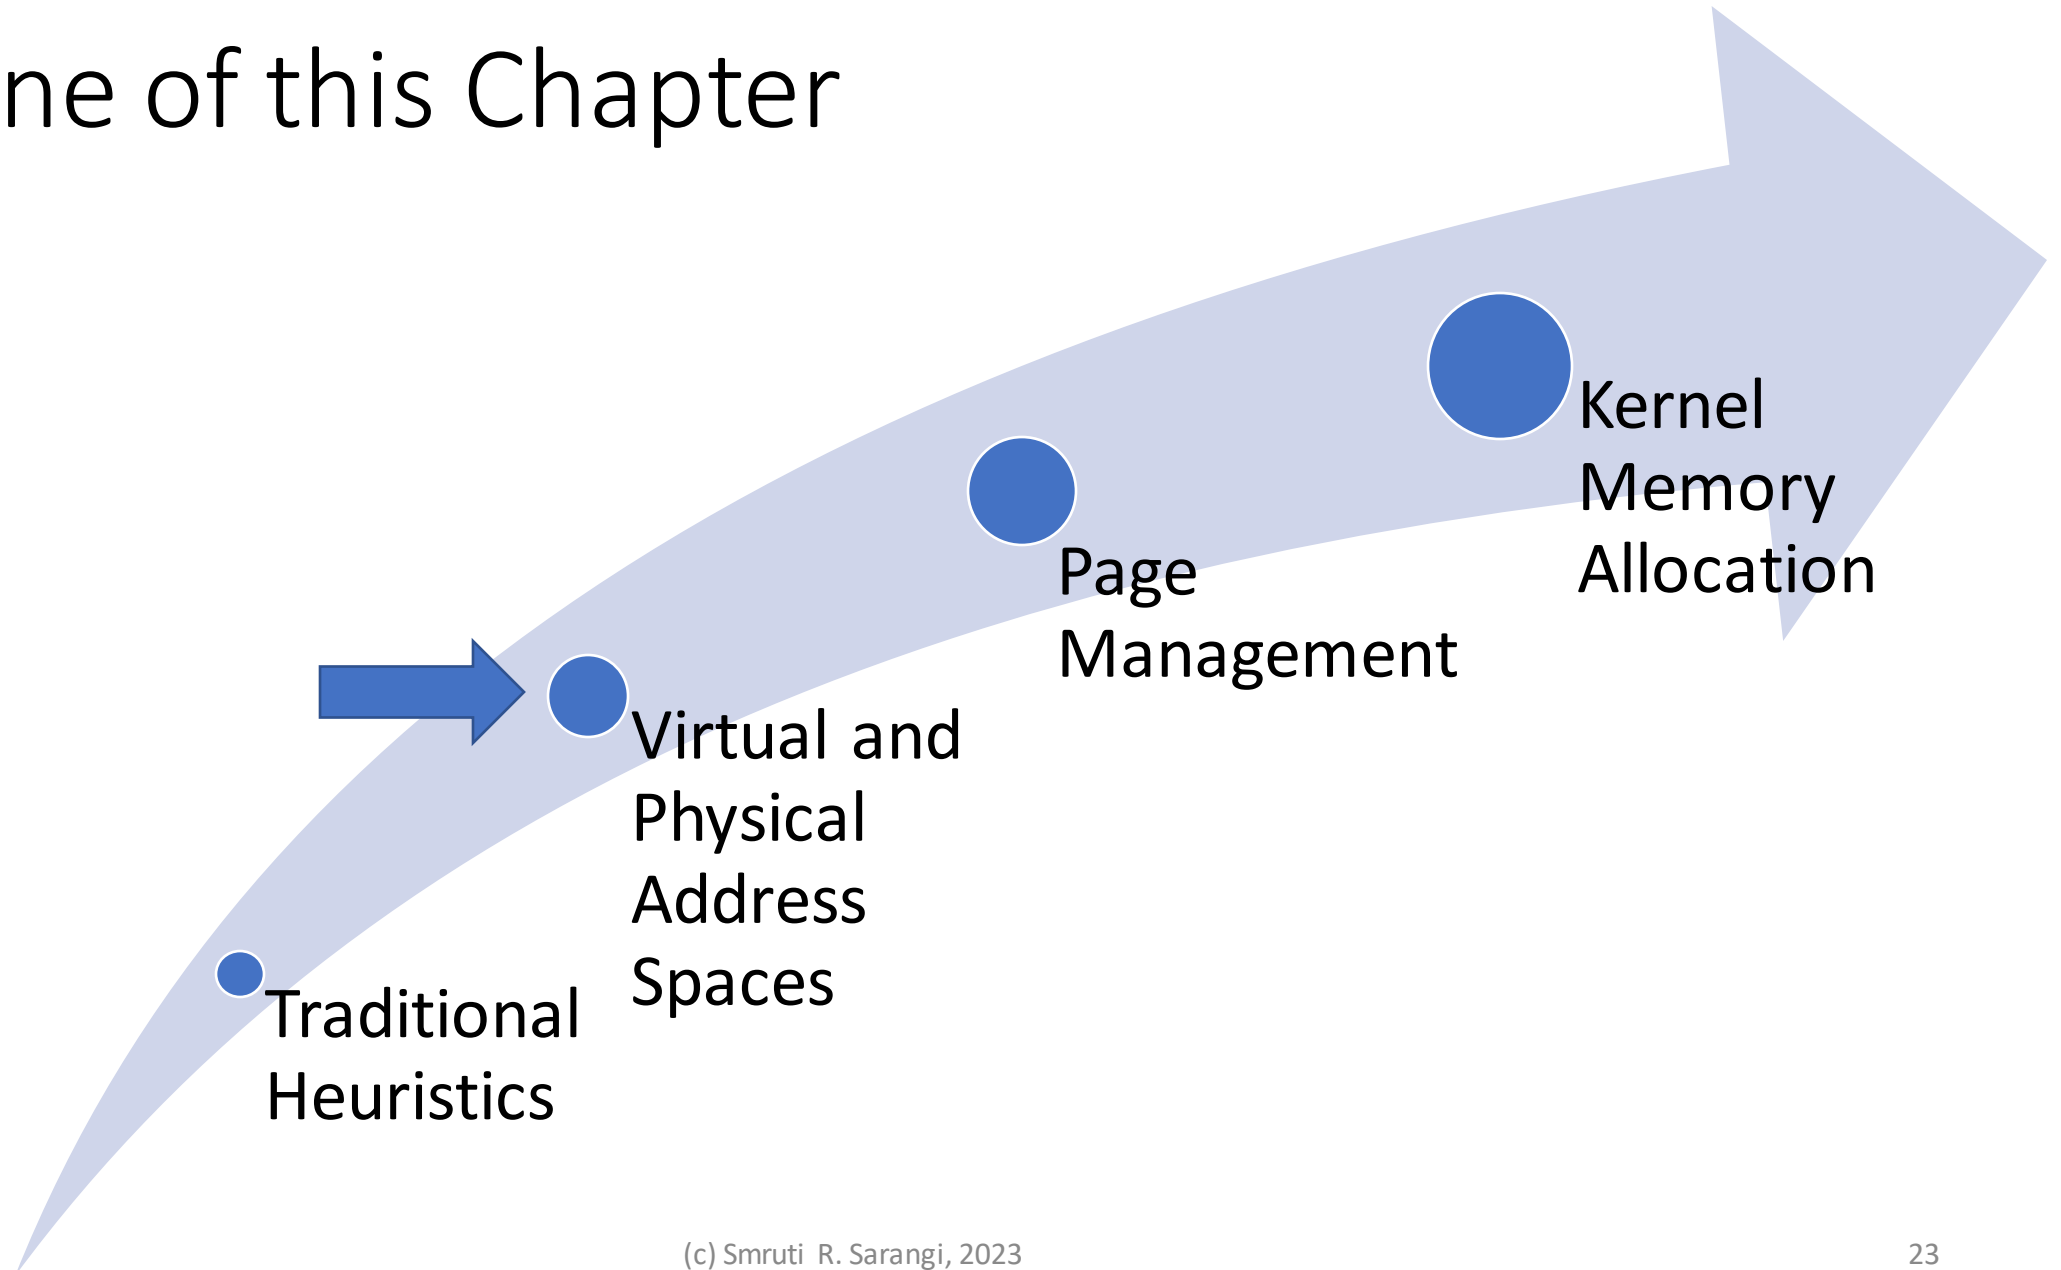


Thrashing

- Consider a **system** with a lot of **processes** and a paucity of memory
- If the space **allocated** to a process is less than its **working** set
 - The process will spend most of its time **fetching** its working set
 - The **performance** counters will indicate **low** CPU utilization
 - The kernel will see the **low** load average and add more **processes** to the CPU's run queue
 - This will make the **problem** even worse
- Ultimately the system will **crash**

This is called thrashing.

Outline of this Chapter




Let us start from *mm_struct*



```
struct mm_struct {  
    ...  
    pgd_t *pgd;  
    ...  
};
```

Pointer to the page table. The CR3 register is set to this value. Type: u64



CR3 register points to the PGD of the current process

Explanation



/arch/x86/include/asm/pgtable_types.h

5-Level Page Tables

| Acronym | Full Form |
|---------|-------------------------|
| PGD | Page Global Directory |
| P4D | Fourth level page table |
| PUD | Page Upper Directory |
| PMD | Page Middle Directory |
| PTE | Page Table Entry |

57-bit virtual address

128 PB VA space

The variable *pgprot_t* contains the protection bits

| Acronym | Full Form |
|------------|----------------------------|
| PROT_READ | Read permission |
| PROT_WRITE | Write permission |
| PROT_EXEC | Execute permission |
| PROT_SEM | Can be used for atomic ops |
| PROT_NONE | Page cannot be accessed |



/include/uapi/asm-generic/mman-common.h

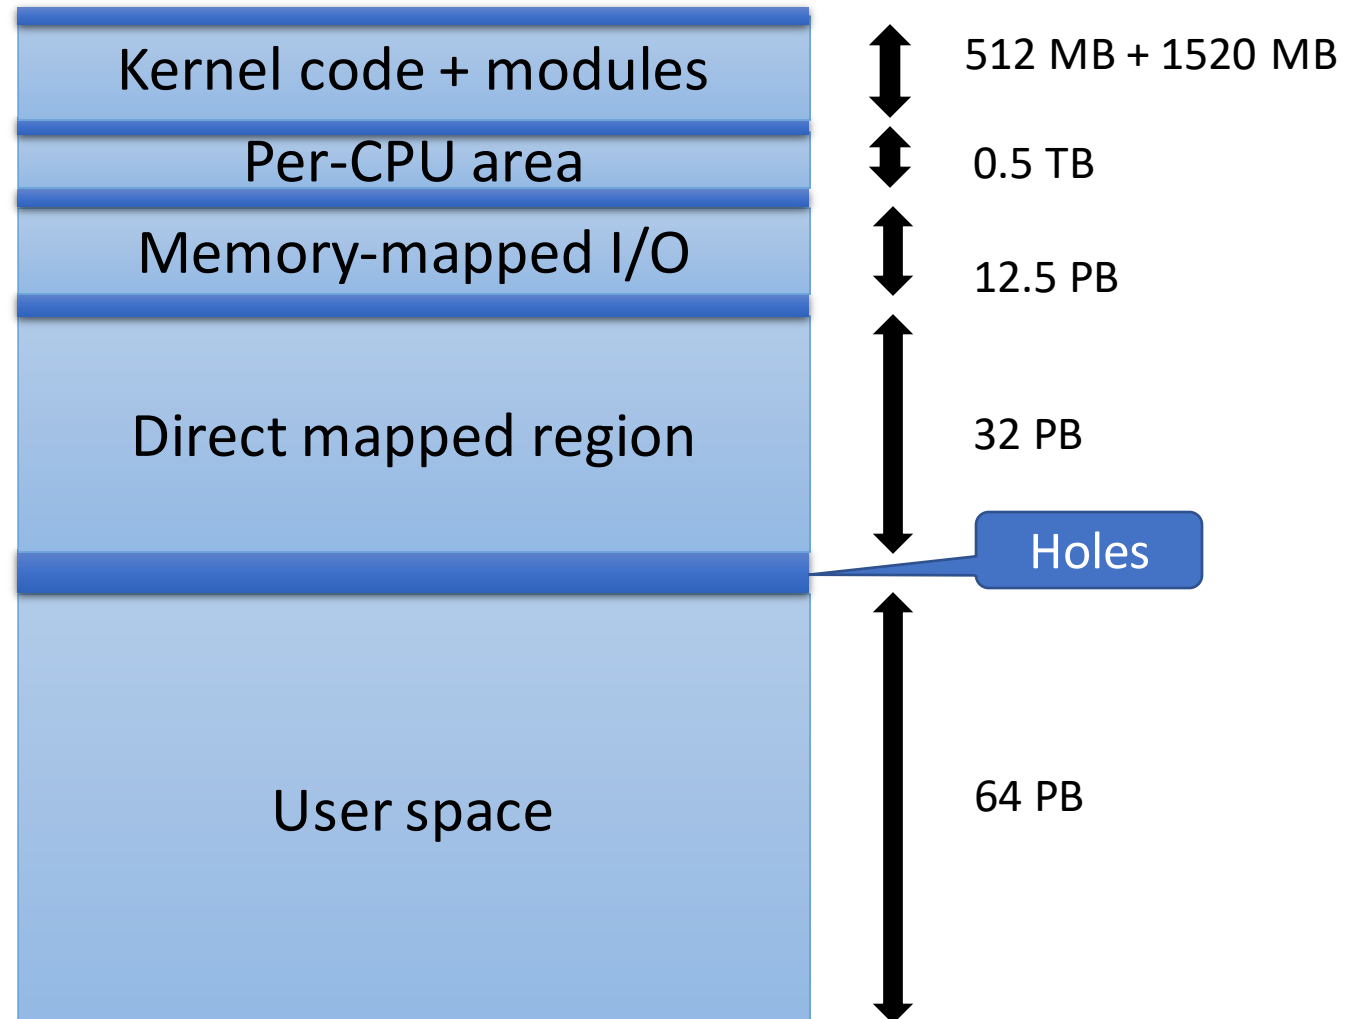
Virtual Memory Map



[Documentation/x86/x86_64/mm.rst](#)



- The **memory** map is not drawn to **scale** and many regions have not been shown.
- There are **holes** between **regions**. If any of the holes are **written** to, it is a **fault**.
- The **direct-mapped** zone can be used to **access** the physical memory directly (albeit by **subtracting** an offset)
- The kernel **text** (code) starts at physical address 0



The *follow_pte* function



/mm/memory.c



Key function used to **traverse** the page table

```
int follow_pte(struct mm_struct *mm, unsigned long address,
               pte_t **ptepp, spinlock_t **ptlp) {
    pgd_t *pgd;
    p4d_t *p4d;
    pud_t *pud;
    pmd_t *pmd;
    pte_t *ptep;

    pgd = pgd_offset(mm, address);
    p4d = p4d_offset(pgd, address);
    pud = pud_offset(p4d, address);
    pmd = pmd_offset(pud, address);
    ptep = pte_offset_map_lock(mm, pmd, address, ptlp);

    *ptepp = ptep;
    return 0;
}
```



This code does not show the **cases** where an **entry** does not exist.

Walk the page table

Set the return value. Keep the page locked.

Accessing any given level (let's say PMD)

```
// /include/linux/pgtable.h
pmd_t *pmd_offset(pud_t *pud, unsigned long address) {
    return pud_pgtable(*pud) + pmd_index(address);
}
```

```
unsigned long pmd_index(unsigned long address) {
    return (address >> PMD_SHIFT) & (PTRS_PER_PMD - 1);
}
```

Extract the index from the address.

```
// /arch/um/include/asm/pgtable-3level.h
#define pud_pgtable(val) ((pmd_t *) __va(pud_val(val) & PAGE_MASK))
```


$\sim(2^{12} - 1)$

```
// /arch/x86/include/asm/page.h
#define __va(x) ((void *)((unsigned long)(x) + PAGE_OFFSET))
```

Starting virtual address of the kernel



Structure of a PTE entry



- Some other **important** data structures that are a part of the **memory** subsystem
 - **page** → Data structure for every **physical** page (frame) in **main memory**. The aim is to record (to some **extent**) who is using the page and what **for**.
 - **folio** → Represents a **contiguous** set of bytes (physically and virtually). Its size is a power of two and is \geq the page size.



struct page



- flags
- One large union (20/40 bytes): it can store a bunch of things. Choices:
 - Pointer to the address space (I/O device whose pages are mapped)
 - Pointer to a pool of pages
 - Page map (mapped from a dedicated device or DMA pages)
- Reference count



This is a classic example of a data structure that has a very flexible structure: it can store anything (depending upon the end user).

<https://lwn.net/Articles/893512/>

<https://lwn.net/Articles/849538/>

struct folio

- **Definition:** A compound page is an **aggregate** of two or more contiguous pages
- A folio primarily **points** to a compound **page**.
 - It is primarily needed to manage millions of pages in large memories
- It points to the **first page** in a **compound** page
- It is very useful for **memory mapped I/O** (I/O devices and files)
- It is naturally aligned towards **read** prefetching and **sequential** writes
- Writes and **modification** bits can be maintained at the **folio** level
 - Easier to maintain LRU-based **replacement** lists

Folio of pages

- A **folio** acts like a single page
- It has its **permission** bits and copy-on-write state
- Whenever a **page** needs to be migrated, swapped out, or replicated (because COW)
 - If a page is a part of a **folio**, then the entire **operation** happens on the folio
- There used to be a concept called **huge pages** (still is)
 - We can have **pages** with size 2 MB and 1 GB
 - Some **server** processors support huge pages. This **requires** changes to addressing or multiple entries are created in the **TLB** (one for each page). The latter solution is very **expensive**.



pte → pfn → page



```
#define pte_pfn(x) phys_to_pfn(x.pte)
#define phys_to_pfn(p) ((p) >> PAGE_SHIFT)
```



```
#define __pfn_to_page(pfn) \
({ unsigned long __pfn = (pfn); \
  struct mem_section *__sec = __pfn_to_section(__pfn); \
  __section_mem_map_addr(__sec) + __pfn; \
})
```

- The **contents** of the *pte* (**page table** entry) contain the physical frame number and other **protection** information
- *pfn* is just the **physical page number** that is obtained by shifting the *pte* by PAGE_SHIFT (=12)
- *struct page* corresponds to a physical frame. Physical pages are **organized** into different **sections**. Each section points to a **memory map** (an array of **page** structures).

TLB

TLB

- It is important to **manage** the TLB well.
- More than **99%** of the requests are satisfied by the TLB.
- A TLB miss is quite **expensive**. It involves a costly page table **walk**.
- In x86 machines
 - By **default**, it is a 4 to 16-way set associative cache
 - Some processors allow the user to **configure** the associativity
 - Some **processors** can also have a 2-level TLB or a separate data TLB and i-TLB
 - Each **entry** of the TLB (corresponding to a **virtual page number**) contains a pointer to a physical page number, and has other **protection** bits (+ other **data**)
- **Flushing** the TLB
 - Just **modifying** the CR3 register **flushes** the TLB
 - Some entries can still be retained (not flushed) if the G (**global**) bit is set
 - The *invlpg* instruction can flush the entire TLB, or a specific page, or the pages belong to a **process**. How ???

Notion of the ASID (PCID)

- A context **switch** is **heavy** on the TLB
- The new **process** suffers from plenty of TLB **misses**
- The same is **true** when the old **process** runs on the **core** again
- If we can **store** the *pid* along with every TLB entry, our job is done
- Then two **processes** can run on the same core (one after the other) without flushing the **TLB**.
- Let us call this additional annotate the ASID (address space ID), Intel calls it PCID → Processor-Context ID

Problem:



A *pid* is an OS-specific concept, whereas the **ASIC/PCID** is a hardware concept. How to **reconcile** both?

Let the Hardware Win

- Maintain a small per-CPU array of PCIDs
- Cache the last few mm (memory maps) for tasks that executed on the CPU
- The PCID bits are set to [1, TLB_NR_DYN_ASIDS]
- TLB_NR_DYN_ASIDS = 6 [default value] (/include/asm/tlbflush.h)
- The PCID is the first 12 bits of the CR3 register
- Each TLB entry also contains the associated PCID
 - It is matched with the current PCID
 - The INVPCID instruction can be used to invalidate all the TLB entries that corresponds to a single PCID




Lazy TLB Mode

<https://notes.shichao.io/utlk/ch2/>

- Assume several CPUs share a **page table**
 - One of them runs a call that **invalidates** an entry
 - This is invalidated for the full **process**
 - All the **TLBs** need to be flushed
- Let's say that one of the CPUs is running a **kernel** thread
 - It need not **invalidate** the entry immediately
 - It can set the **CPU** to "**lazy** TLB mode"
 - The kernel threads don't have **separate** page tables. It is a common one that is **appended** with user mode page tables (one large user + kernel page table)


Lazy TLB Mode

Three cases

-  The kernel tries to **access** the **invalidated** page. This will happen only via fixed entry points. This cannot happen in an uncoordinated fashion. Appropriate checks can be carried out and exceptions can be thrown.
-  The kernel **switches** to another user process. In this case, all the TLB entries of the original process are **flushed** out.
-  The kernel switches back to the **same** user process. Just finish the work of **invalidating** all the entries that were **deferred**.

Partitioning Physical Memory

NUMA Machine



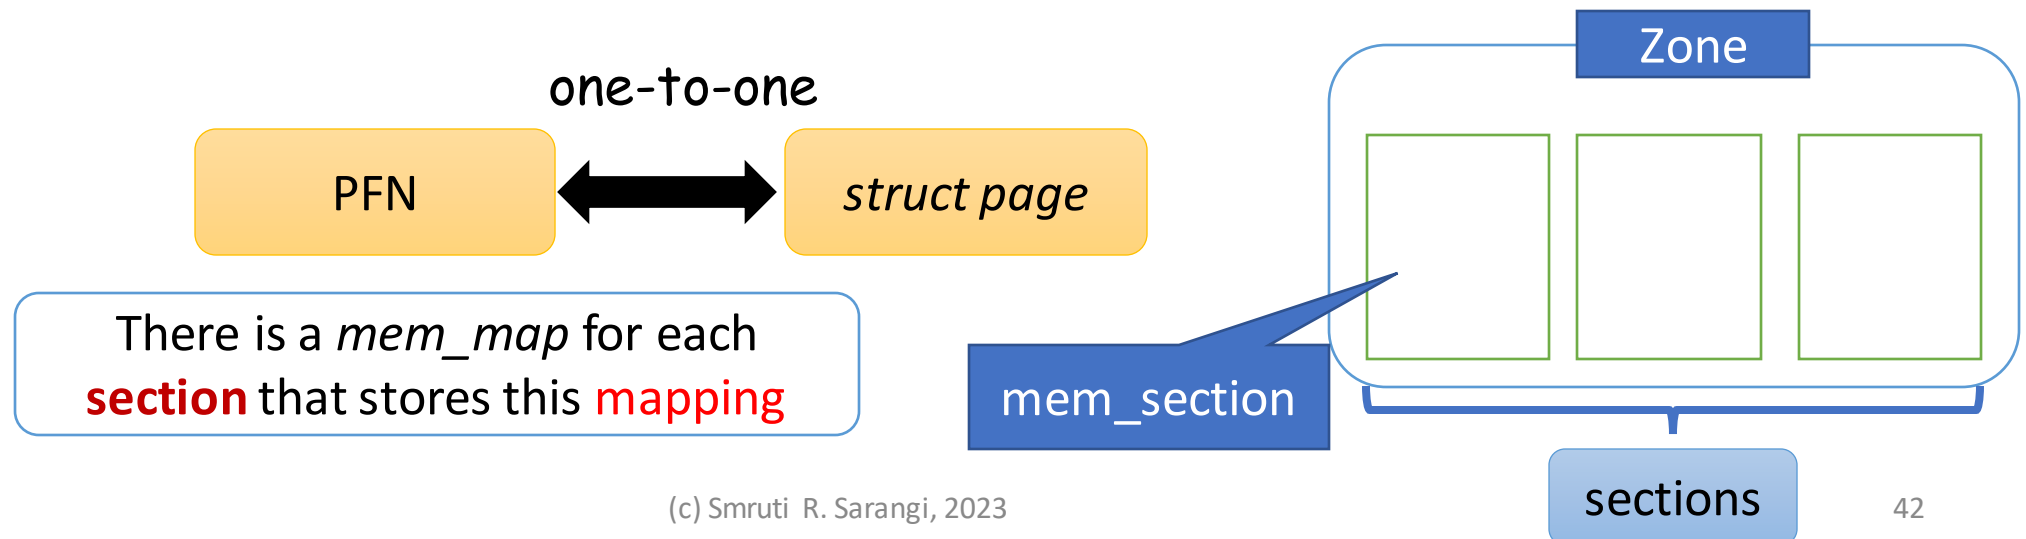
- CPUs have some amount of local memory, which is much faster
- They are also organized into clusters.
- A CPU can access all memory: intra-cluster and inter-cluster
- Accesses to intra-cluster memory is much faster.
- This is a non-uniform memory access machine (NUMA machine)
- Each cluster is known as a node

Is all physical memory in a node the same?

No!

<https://lwn.net/Articles/789304/>

1. Partition the physical memory space (in a **node**) into **zones**
2. Treat the **frames** (**physical** pages) in each zone **differently**.
3. The **zones** may themselves be **stored** on different devices.
4. I/O Devices may need a **dedicated** region (assign a zone)



Physical Memory Zones in Linux



[/include/linux/mmzone.h](#)

```
enum zone_type {  
    ZONE_DMA,  
    ZONE_NORMAL,  
#ifdef CONFIG_HIGHMEM  
    ZONE_HIGHMEM,  
#endif  
    ZONE_MOVABLE,  
#ifdef CONFIG_ZONE_DEVICE  
    ZONE_DEVICE,  
#endif  
    __MAX_NR_ZONES  
};
```

Physical pages that are accessible only by the DMA controller

Normal frames

Useful in systems where the physical memory exceeds the size of max virtual memory.

It is assumed that the corresponding memory device may be removed at any point of time and possibly re-inserted later.

These frames are stored in novel memory devices like NVM devices.

struct zone



```
struct zone {  
    int node;  
  
    struct pglist_data *zone_pgdat;  
    unsigned long zone_start_pfn;  
  
    atomic_long_t managed_pages;  
    unsigned long spanned_pages;  
    unsigned long present_pages;  
  
    const char *name;  
  
    struct free_area free_area[MAX_ORDER];  
}
```

Pointer to the *pglist_data* structure

$\text{zone_end_pfn} = \text{zone_start_pfn} + \text{spanned_pages} - 1$

Name of the zone

Free areas in a zone
(physical memory)

Data Structure to Manage the Memory in a Zone



```
typedef struct pglist_data {  
    struct zone node_zones[MAX_NR_ZONES];  
    struct zonelist node_zonelists[MAX_ZONELISTS];  
    int nr_zones;  
  
    unsigned long node_present_pages;  
    unsigned long node_spanned_pages;  
    int node_id;  
  
    struct task_struct *kswapd;  
    struct lruvec __lruvec;  
} pg_data_t;
```

Ordering of zones

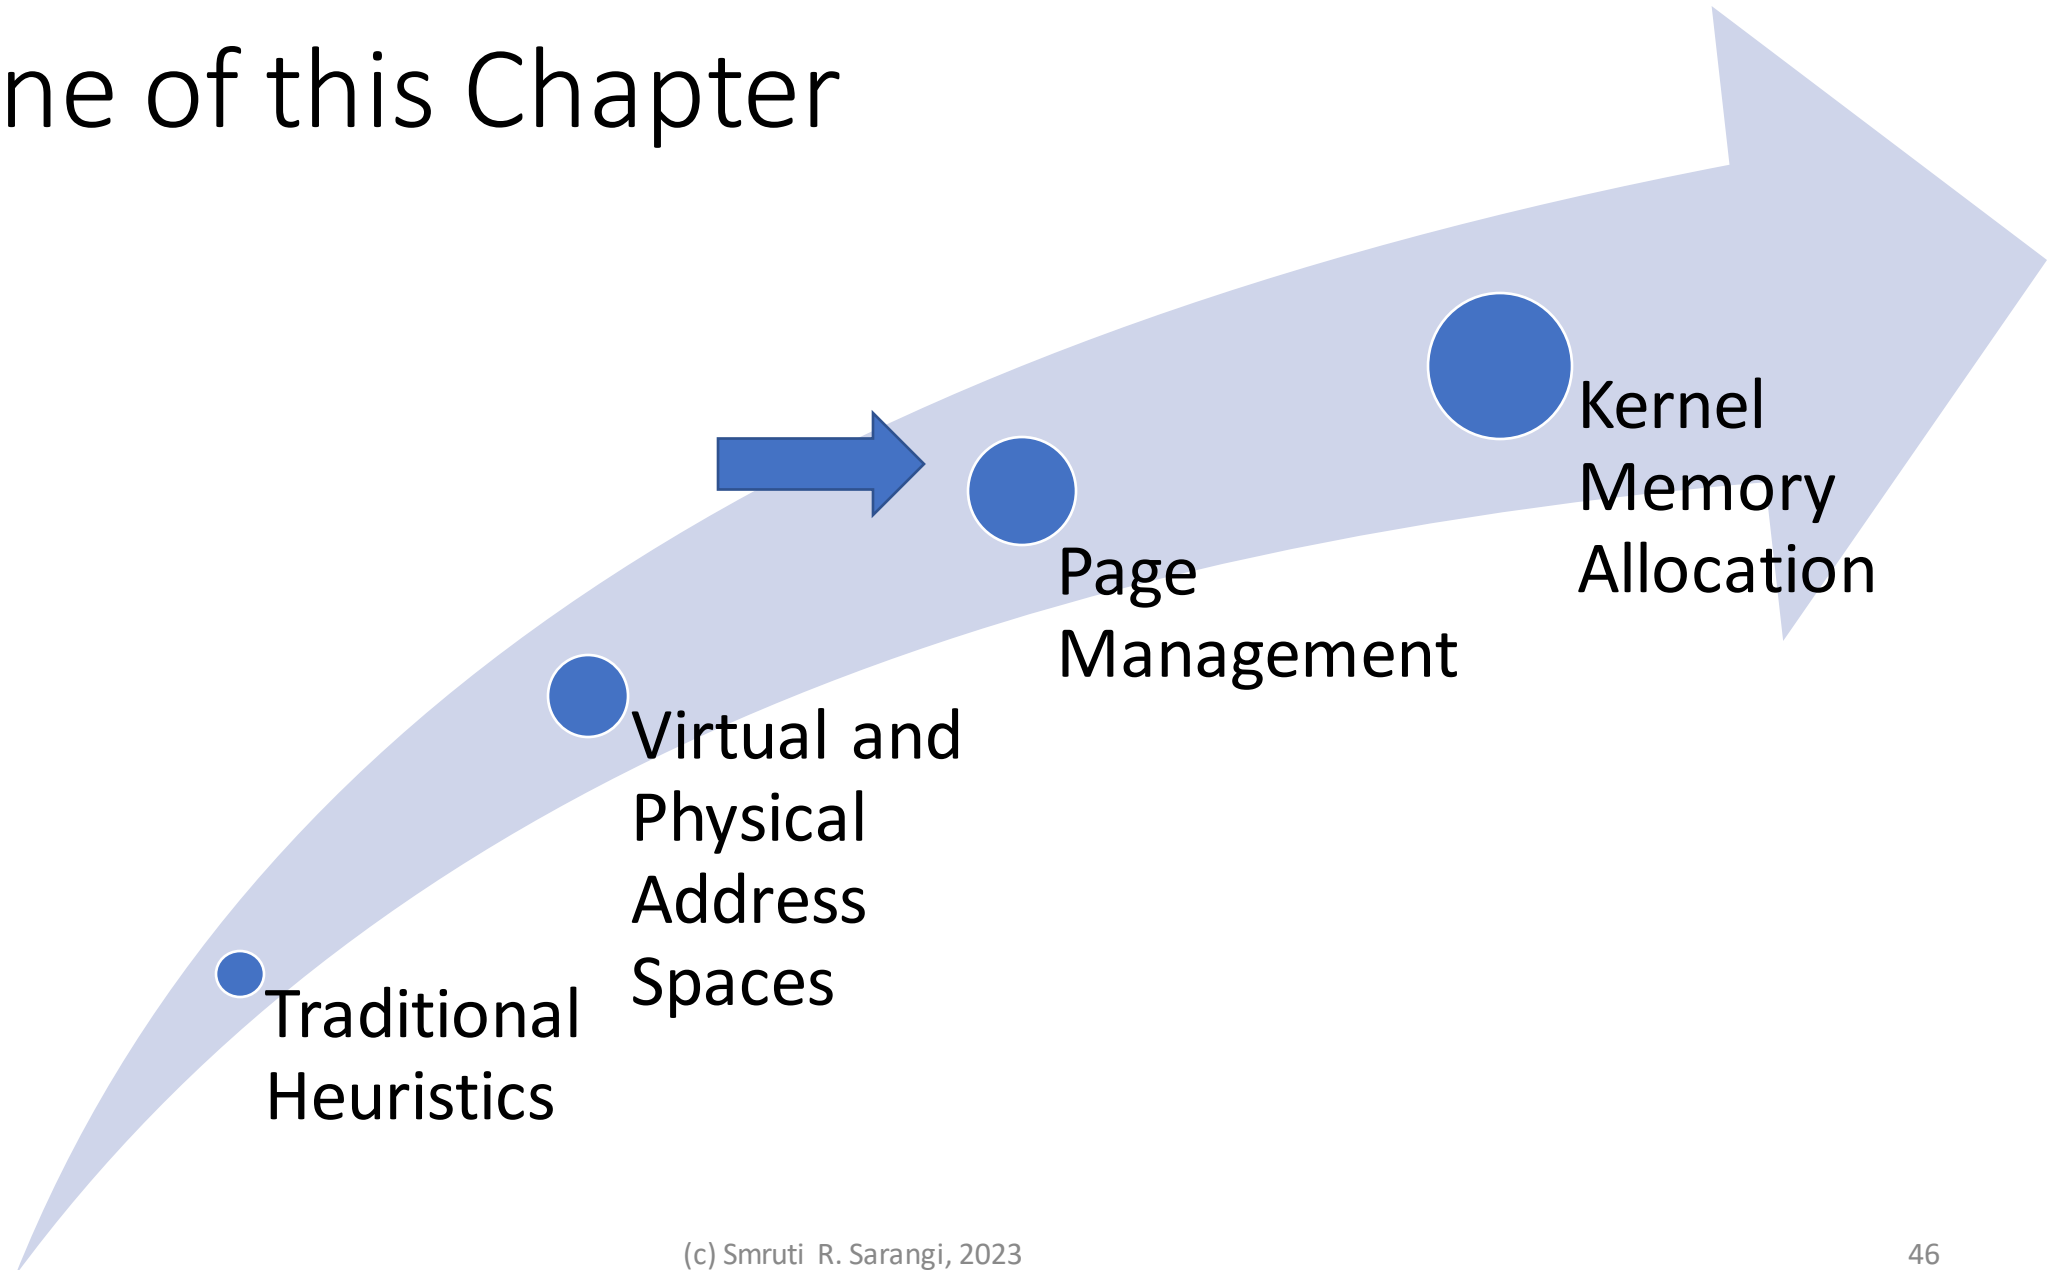
Zones organized hierarchically

Number of pages owned by NUMA node_id

Page swapping daemon

Maintains LRU state information

Outline of this Chapter





Required Background

Bloom Filters

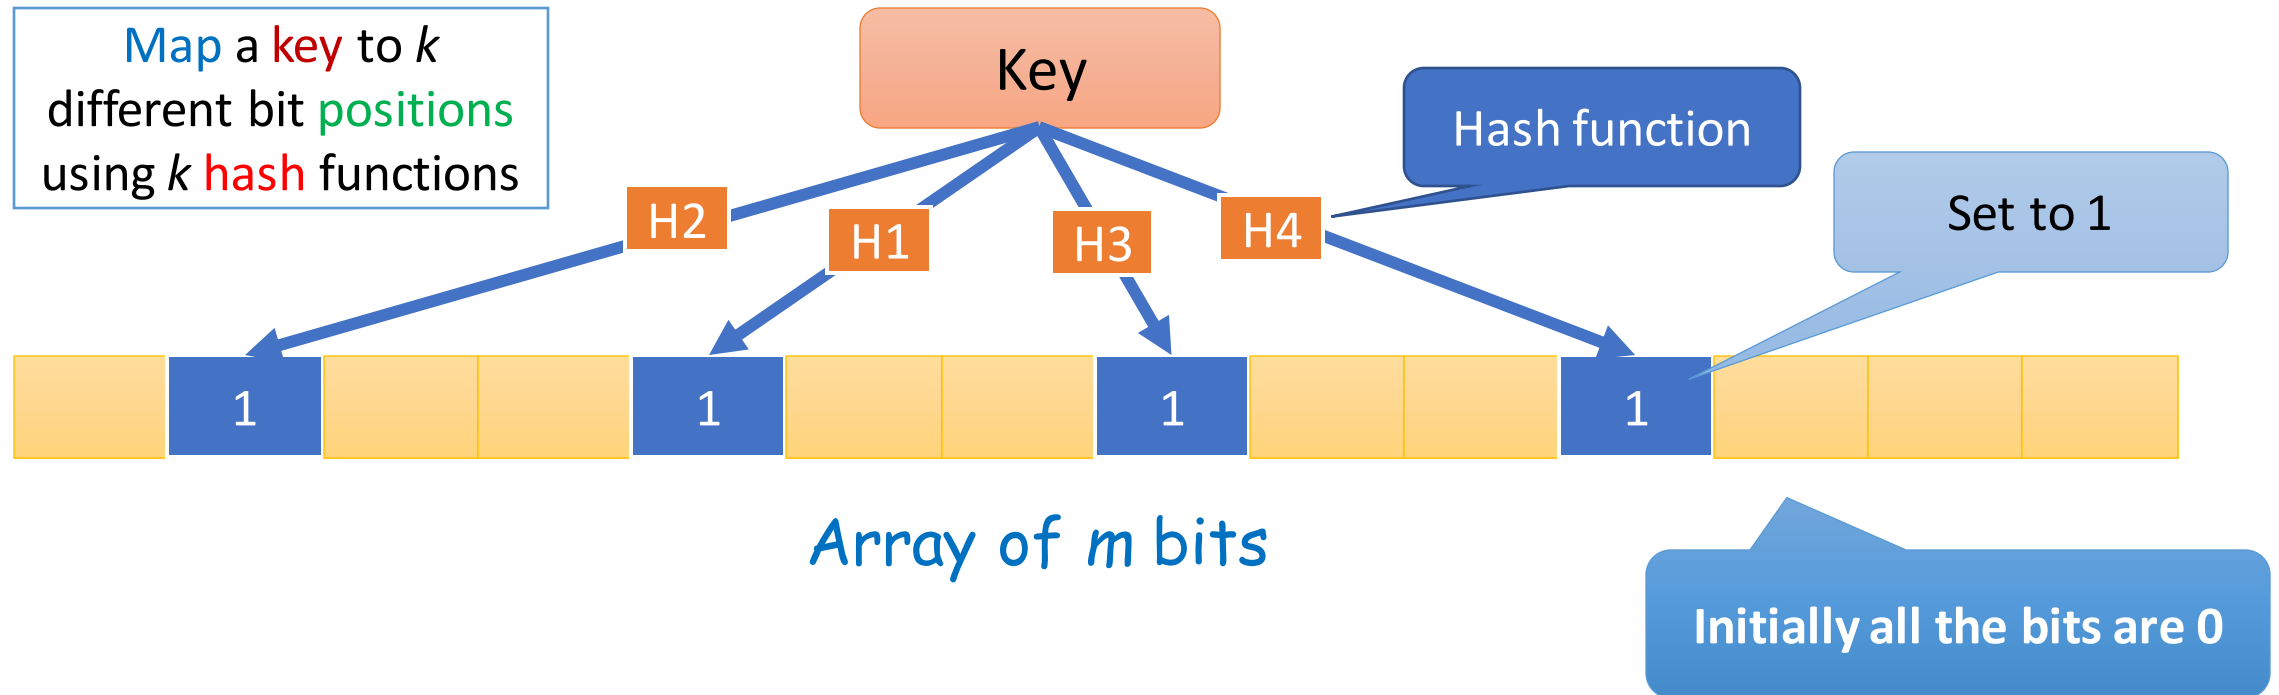
PI Controller

Reverse Maps

Inserting a Key in a Bloom Filter



A **data structure** that answers if an **element** is a member of a set (**probabilistically**)



Checking for Set Membership in a Bloom Filter

1. Given a key **compute** the k different **hash** values (bit positions)
2. Check that each bit **position** stores a 1. If at least one bit position stores a 0, then the key is not **present**.
3. Elements cannot be **deleted** unless we store a **count** at each **bit** position
4. We need to periodically **reset** the Bloom Filter



False positives are allowed

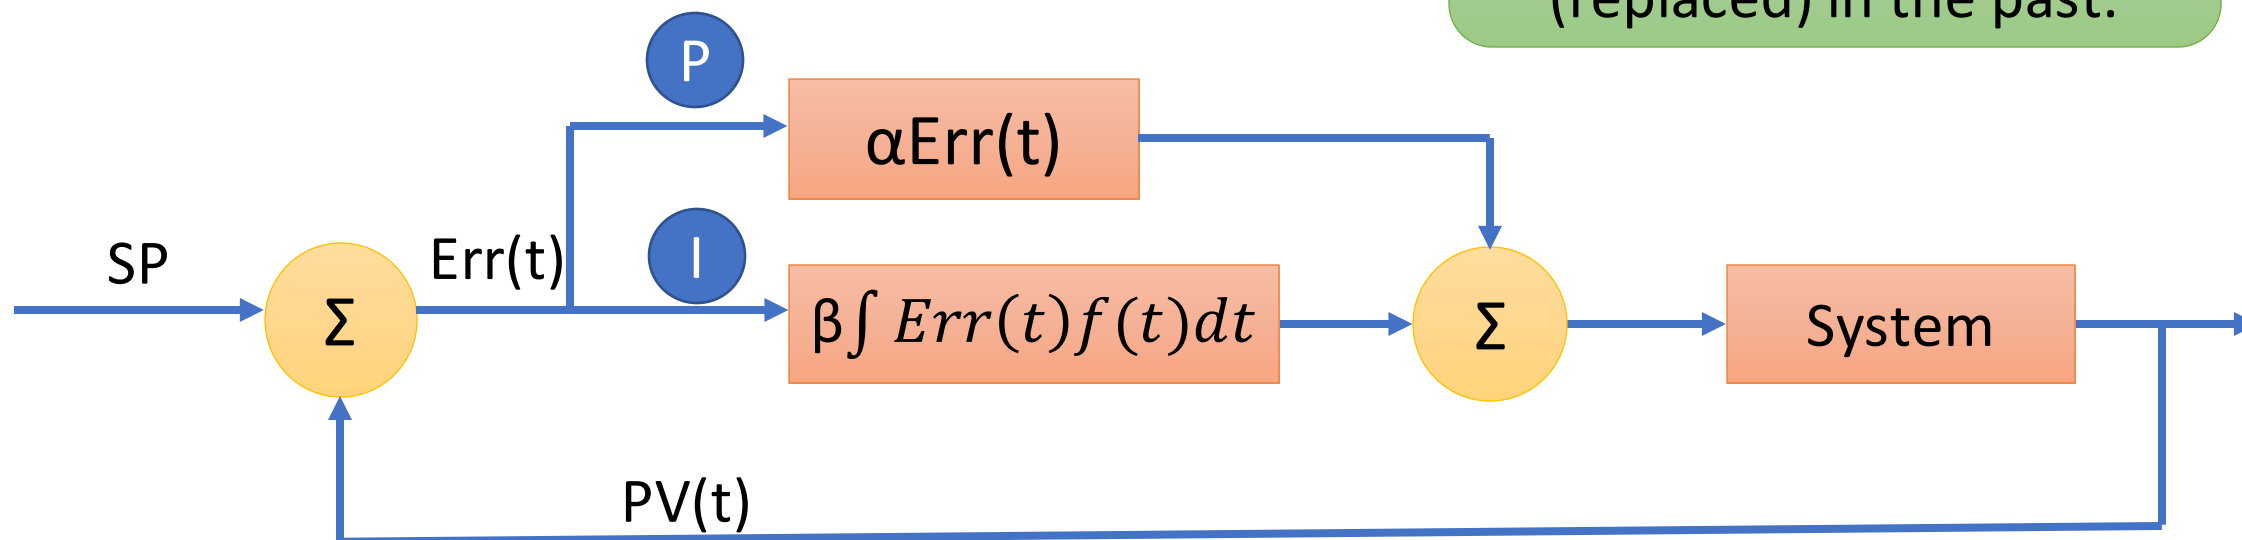


No false negatives

Background

The PI Controller

A refault is a page fault for a page that was evicted (replaced) in the past.



| Term | Meaning | Explanation in this context |
|--------|-------------------|--|
| P term | Proportional part | refaulted/ (evicted + pinned) |
| I term | Integral part | Exponential moving average of P (factor = ½) |
| SP | Set point | Refault rate of the first tier of ANON |
| PV | Process variable | Refault rate of the first tier of FILE |

Both should be the same (target)

Error Term



```
err = (pv->refaulted < 64 ) ||  
      pv->refaulted /pv->total * swappiness <=  
      (sp->refaulted + 1) / (sp->total + 64)* (200 - swappiness)
```

err = 1 *if*

1. Too few pages have **refaulted** (< 64)
2. The refault rate (refaulted/total) is **lower** than the **set point** (subject to the **swappiness**)
3. The swappiness is close to 0 (**low**)

Reverse Mapping (called from *shrink_folio_list*)

Reverse map (rmap)

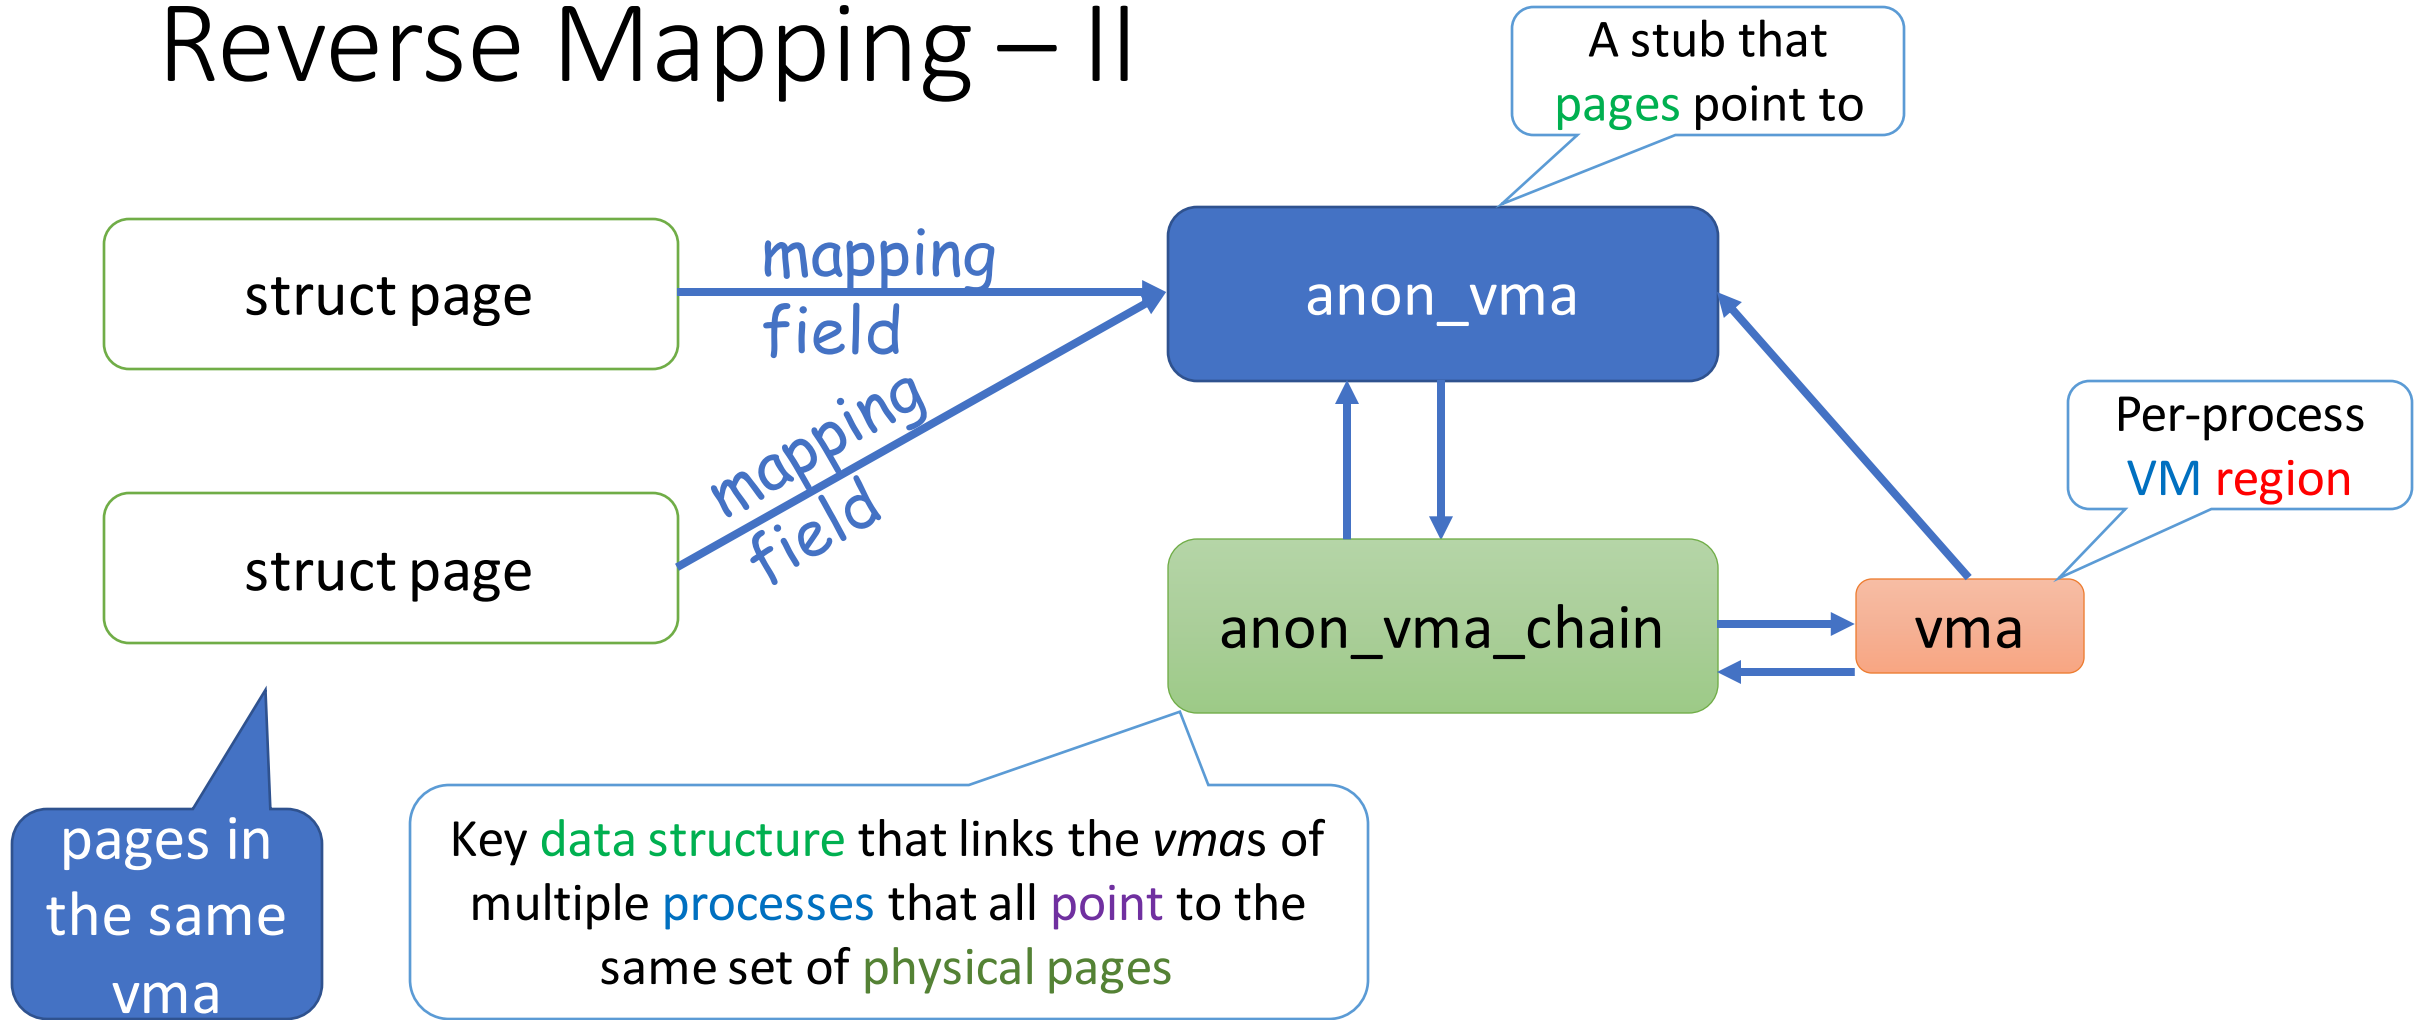
struct page



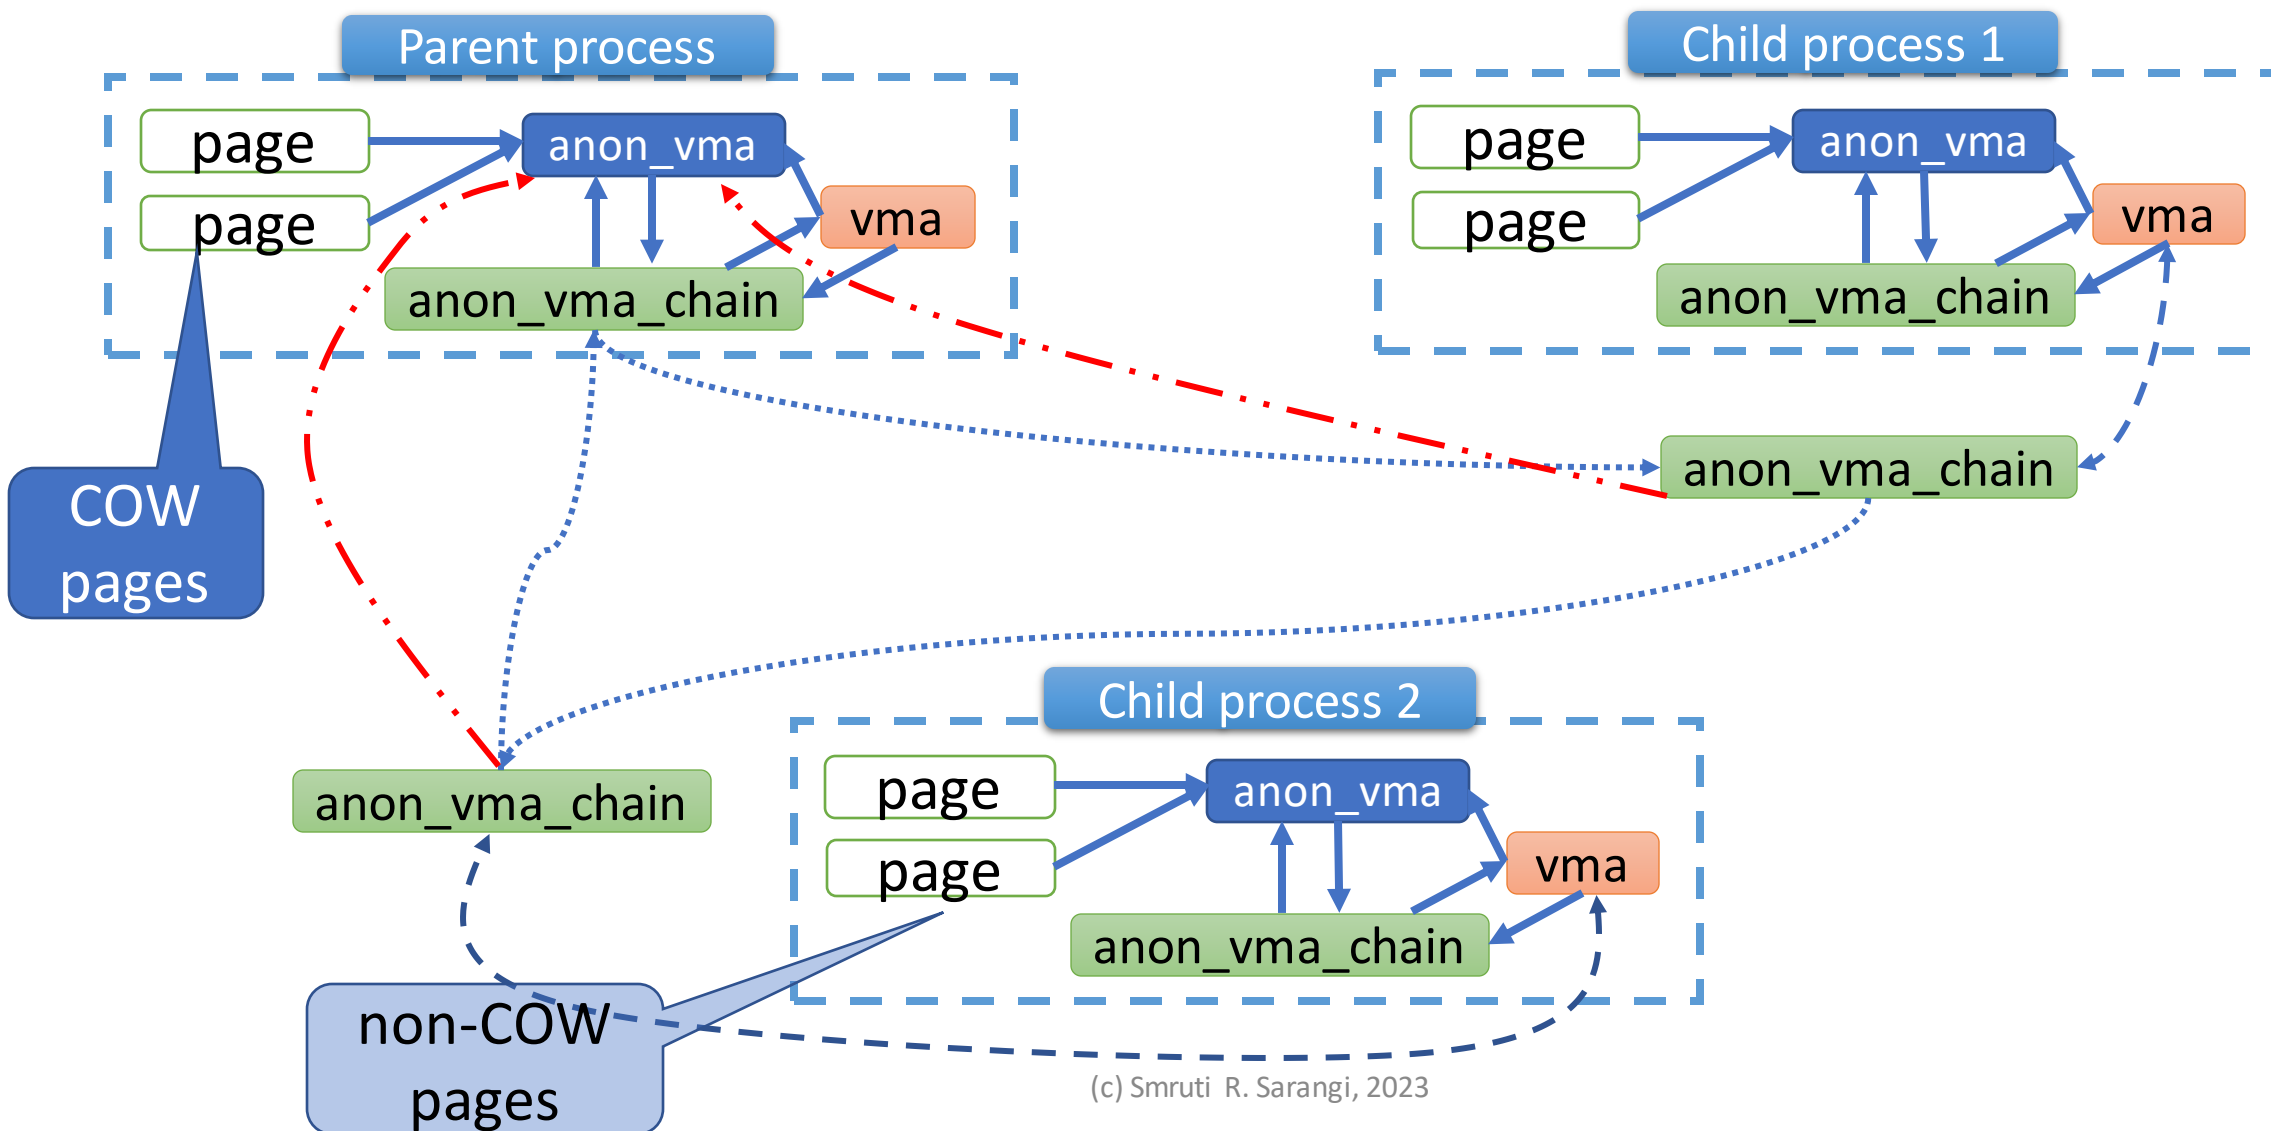
List of PTEs

- While **freeing** a page (or **folio**), we need to **consider** the possibility of it being mapped to **multiple** address spaces (across processes)
- We would thus need to create a reverse map (rmap)
 - Map a **struct page** (or **folio**) to **PTEs** across processes
 - **Swapping** out a folio requires a **rmap** walk
 - Each PTE entry needs to be **processed**
- The idea is to **efficiently** associate a list of *vma regions* with each **PTE**

Reverse Mapping – II




Multi-process View



anon_vma_fork in *mm/rmap.c*

- **Justification** of the complex structure →
- The rest of the code **relies** on this structure



What is the **need** for the **extra** *anon_vma_chain* associated with a **child** process?

- An *anon_vma_chain* is associated with a **single** *anon_vma*
- We need one to **point** to the parent process's pages (via the **corresponding** *anon_vma*) that are also mapped to the child (COW pages). This is the job of the extra *anon_vma_chain*.
- The other one is needed for child pages in a *vma* to point to. These pages **exclusively** belong to the child.

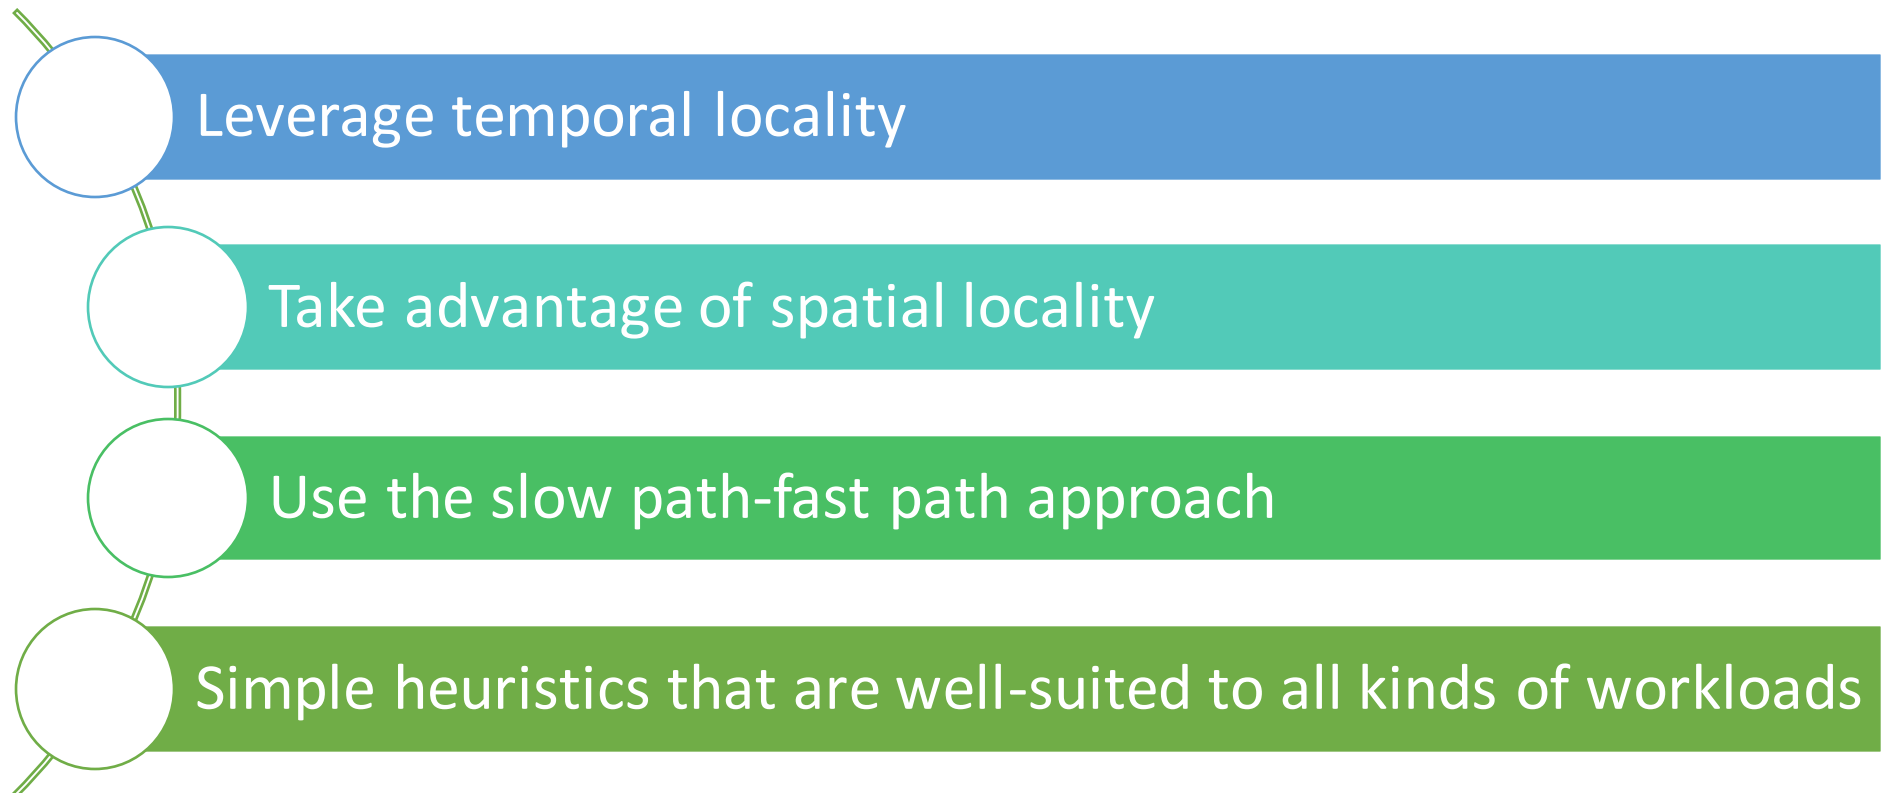


Go forward

Page Replacement

New MG-LRU page replacement algorithm

- **Objectives** of a page replacement algorithm.





```
struct lruvec {  
    struct pglist_data *pgdat;  
    unsigned long refaults [ANON_AND_FILE];  
    struct lru_gen_struct      lrugen;  
    struct lru_gen_mm_state    mm_state;  
};
```

Number of refaults



```
struct lru_gen_struct {  
    unsigned long max_seq;  
    unsigned long min_seq[ANON_AND_FILE];  
    unsigned long timestamps[MAX_NR_GENS];  
    struct list_head  
        lists[MAX_NR_GENS][ANON_AND_FILE][MAX_NR_ZONES];  
};
```

age at least *min_ttl* ms
after *min_seq* is created

Latest

Oldest: separate for anon
and file types

lru_gen_mm_state




```
struct lru_gen_mm_state {
    unsigned long seq;
    struct list_head *head;
    struct list_head *tail;
    unsigned long *filters[NR_BLOOM_FILTERS];
    int nr_walkers;
};
```

Current sequence
number

tail has finished.
continue at *head*

list of Bloom
filters

#page walkers



Array of *mm_struct*
objects



The basic idea of compressing the memory footprint

- Iterate through the list of **zones**
- **Shrink** each zone

<https://lwn.net/Articles/419713/>



Multi-Gen LRU Algorithm

- A **group** of **pages** is associated with a generation.
- The generation indicates **recency of access**.
- Take the **priority** of pages into account. Prefer **unmapped** and **clean** pages.
- **Evict** pages in the **oldest** generation (min gen) and also take the **priority** into account



Tracking Accesses



Needed to maintain generations

```
pte_t pte_mkold(pte_t pte)
{
    return pte_clear_flags(pte, _PAGE_ACCESSED);
}
```

- This **function** is called by a bunch of functions that **walk** the page table
- The **key idea** is to clear the **_PAGE_ACCESSED** bit in the PTE entry
- The next time that the HW **accesses** this page
 - It can either **raise** a page fault
 - OR **set** the bit itself (**directly**)
- A second scan of this memory region will indicate, which **pages** have been **accessed**.

Indicator of recency

When do we age and subsequently evict?

kswapd

- This is a **daemon** that runs **periodically**: ages and reclaims pages

Call to page allocation routines



Should we run the Aging Algorithm?

should_run_aging(...)



- **Youngest** generation number: `lrugen->maxseq`
- **Oldest** generation number: `lrugen->minseq[ANON]` and `lrugen->minseq[FILE]`
- When $\text{max_seq} - \text{min_seq} + 1 \leq \text{MIN_NR_GENS}$, increment **max_seq**
- Given that `max_seq` and `min_seq` increase **monotonically**, not incrementing the sequence of a **folio** automatically **ages** it
- A **sliding window** of generations is **maintained**. It tracks `[MIN_NR_GENS, MAX_NR_GENS]` generations



Triggered by a call to age the *lruec* and reclaim pages

The Aging Process of Pages

Corner
case

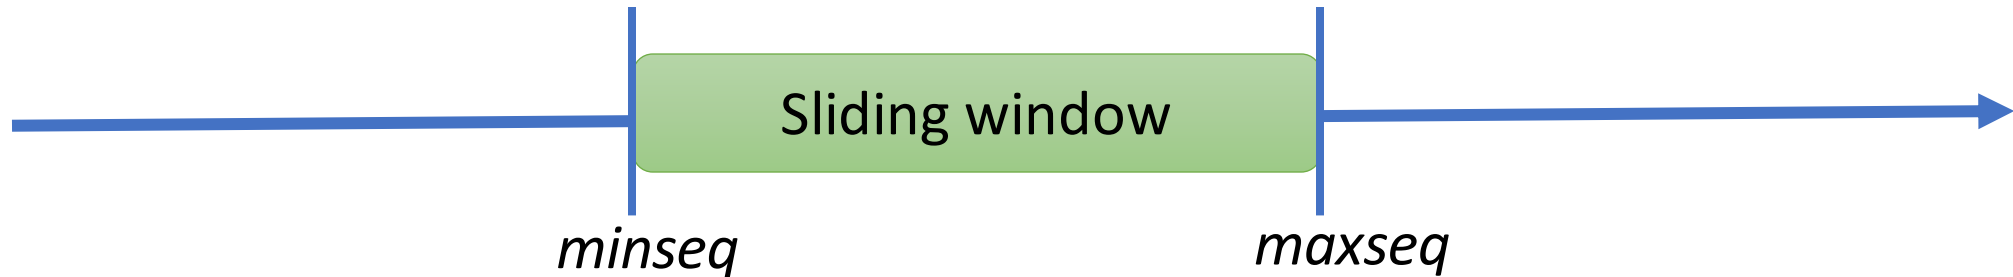


$\text{max_seq} - \text{min_seq} == \text{MIN_NR_GENS}$

```
if (young * MIN_NR_GENS > total)
    return true;
if (old * (MIN_NR_GENS + 2) < total)
    return true;
return false;
```

Too many young pages.

Too few old pages.



Note the
definition

In this function, young page's $\text{gen} == \text{max_seq}$
old page's $\text{gen} == \text{min_seq}$

The Aging Process Triggers *walk_mm*

Walk through all the `mm_structs`

- Walk through the individual page tables

Skip a PMD's pages if

Eviction Algorithm

Overview

- When the `lrugen->lists[min_seq][type][]` is empty, *min_seq* can be incremented for the corresponding type (*anon* or *file*). Otherwise, **incrementing** \Rightarrow **eviction**
- **Maintain** tiers for each generation – maintained in `folio->flags`
 - **tier** = \log_2 (#recorded_page_accesses)
 - **Tracked** via system calls for file page accesses
 - **Tracked** via access bit changes for ANON page accesses
 - The **first** tier has **single-use clean pages** that are **unmapped**
- **Broad idea**: Compare the *file* and *anon min_seq*[] values
 - Choose the **one** with the lower value
 - If both are the **same**, then **choose that type** whose **first** tier has a lower normalized **refault rate**

Typically in the page cache: SW cache for pages read from the disk



Evict the oldest generation

Main **function** called from *shrink_zones*

evict_folios (lruvec, int swappiness,)

Hyperparameter, typical value: 60

1

scanned = isolate_folios (lruvec, swappiness, ...)

Identify the folios that can possibly be **evicted**.

2

Reclaim pages by calling *shrink_folio_list*

Try to **remove** the folios that have been **identified**. It is possible that a folio is mapped to **multiple** processes.

Do additional bookkeeping and account for all corner cases.



evict_folios → *isolate_folios*



/mm/vmscan.c

Find the type that
should be evicted

type

ANON = 0, FILE = 1, ANON_AND_FILE = 2



```
if (!swappiness)
    type = LRU_GEN_FILE;
else if (min_seq[LRU_GEN_ANON] < min_seq[LRU_GEN_FILE])
    type = LRU_GEN_ANON;
else if (swappiness == 1)
    type = LRU_GEN_FILE;
else if (swappiness == 200)
    type = LRU_GEN_ANON;
else
    type = get_type_to_scan(lruvec, swappiness, &tier);
```

Default if **swappiness** = 0

ANON is **older**

Choose FILE if **swappiness** = 1

Choose ANON if **swappiness** = 200



*get_type_to_scan (lruvec, swappiness, int *tier_idx)*

Collect the number of total and refault accesses



Choose the type that has the lower normalized refault rate

Set point is ANON, tier = 0



```
int gain[ANON_AND_FILE] = {swappiness, 200 - swappiness};
```

```
read_ctrl_pos(lruvec, LRU_GEN_ANON, 0, gain[LRU_GEN_ANON], &sp);  
read_ctrl_pos(lruvec, LRU_GEN_FILE, 0, gain[LRU_GEN_FILE], &pv);  
type = positive_ctrl_err(&sp, &pv);
```

```
return type; /* If (Error = 1) return FILE else ANON */
```

PI
cntrl

logic for comparing the normalized refault rate

The process variable:
refault rate of tier=0
for FILE

The rest of the *isolate_folios* function




```
for (i = !swappiness; i < ANON_AND_FILE; i++) {  
    if(scanned = scan_folios(lruvec, type, list, ...) break;  
  
    /* reset the type*/  
    return scanned;  
}
```

Check if there are
evictable folios and
return the number of
such folios

more about scan_folios




evict_folios → *isolate_folios* → *scan_folios*



scanned

list



scanned

list

1. **Iterate** through the **zones** that need to be **shrunk** (order: **descending**)
 - I. Iterate through all the **folios** pointed to by
lrugen->lists[min_seq][type][zone]
 - a) Check that the **folio** can be evicted
 - b) If the **folio** is pinned, is being **written** back, was recently accessed, or there is a **race** condition, **skip** it
2. Return the number of pages that can possibly be evicted, and a list of folios
 - I. This is the variable **scanned**
 - II. list** variable → list of possibly **evictable** folios

More about *scan_folios*

Does some additional computation

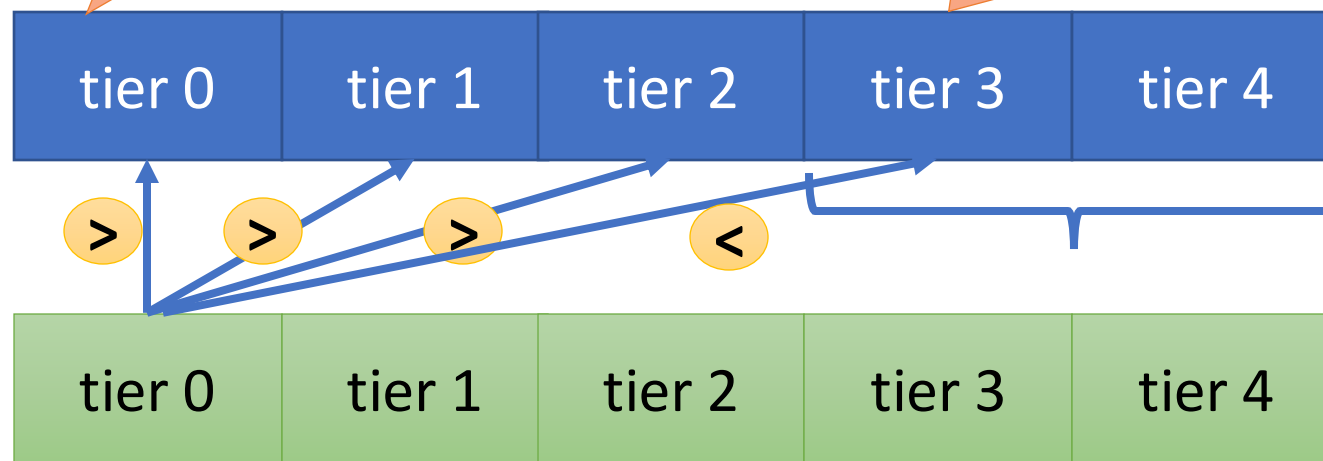
Type chosen for eviction (FILE/ANON)

Other type

lower normalized
refault rate

tier at which the
tier's refault rate is
higher

Increment the
generations of
folios in these
tiers when
they are
scanned



Even if the generation is *min_seq*, if the folio is in a higher tier, give it another chance by incrementing its generation



evict_folios → *shrink_folio_list*

We have a list of folios that can possibly be evicted



1. Perform basic **bookkeeping**. Take the number of **reclaimed** pages into account.
2. **Folios** that are in the process of being written back may **stall** this process. **Wait**.
3. Given a **folio** that needs to be **evicted**
 - i. Look around **nearby** addresses: *lru_gen_look_around* →
 - ii. Mark a few **old** (clear the **access** bit)
 - iii. Note down **PMD** entries that point to primarily **young** (recently accessed) pages (in the Bloom filter)
4. Split **large** folios into smaller folios (for **performance** reasons)
5. **Free** buffers, **flush** the corresponding entries from the **TLB**, **write back**

lru_gen_look_around: main input vm_area

1. Consider the **memory** space that is mapped by one entry in the **PMD**
2. Further limit the **size** of the considered **region** to 64 pages
3. Iterate through all the **pages**
 - I. If it is an **old** page (not recently accessed) continue.
 - II. Get the **folio** associated with the **page**
 - III. Mark the page **old** (**clear** its access bit)
 - IV. If the folio's generation \neq max_seq, **record** the page number in a bitmap
4. If the **number** of **young** pages is more than a certain **threshold** record the PMD number in the **Bloom filter** (**makes page walking more efficient**)
5. Set the **folios** of the **pages** (**recorded** in the bitmap) to the latest generation (**max_seq**) and do the rest of the **bookkeeping**.



? How do we use the Bloom filter (subsequently)?

1. *kswapd*, sequential write commands, and page allocation modules try to **age** and **evict** pages.
2. They call the function *walk_mm()* that walks the page table and marks **young (accessed)** entries as **old (unaccessed)**, as well as updates the **generation** of young entries.
3. Entries that are already old are **skipped**.
4. While walking the PMD **tables**, test if the PMD address is in the **Bloom** filter. If it is not there, then the **PMD** range is most likely **old** and should be **skipped**. Improves **performance**.

Miscellaneous Topics: Thrashing

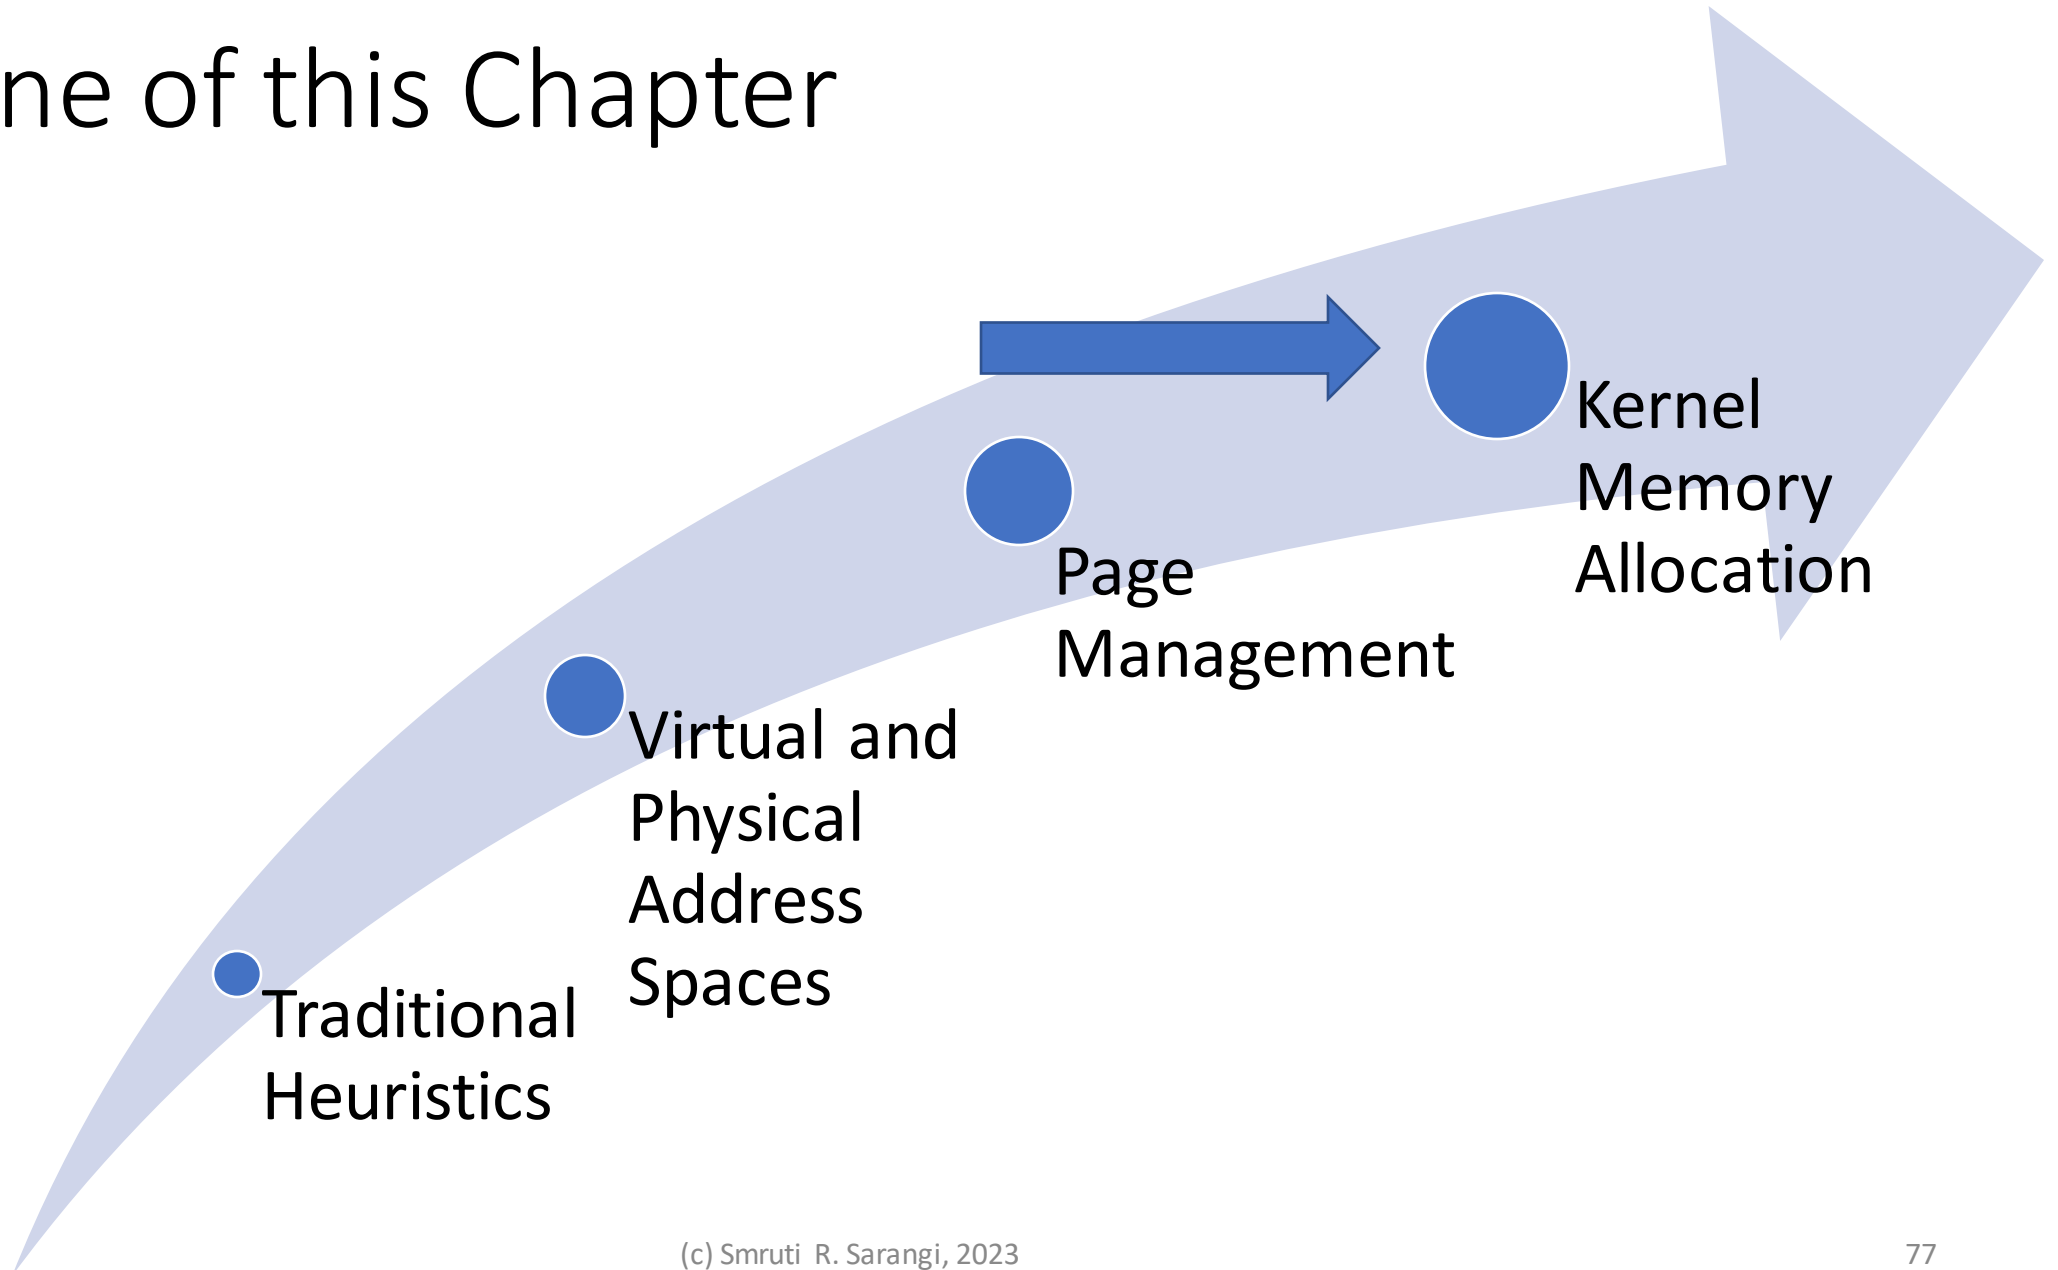
- There is a need to keep the **working** set in memory
- Don't **evict** a page that is less than N ms **old** in **memory**
- $N = 1000$ is practically a **good value** for GUI-based applications
- If the working set cannot be **maintained** in memory
 - The **OOM** killer (out-of-memory) comes into action and **terminates** applications
- The **behavior** is configurable (**hyperparameters**)

MIN_NR_GENS

MAX_NR_GENS

swappiness

Outline of this Chapter



Buddy Allocation

- To manage the physical memory, the **buddy allocation** technique is used in the kernel
- It can provide **physical memory** chunks of arbitrary sizes to **processes**
- It now has more **relevance** given the prevalence of **folios**

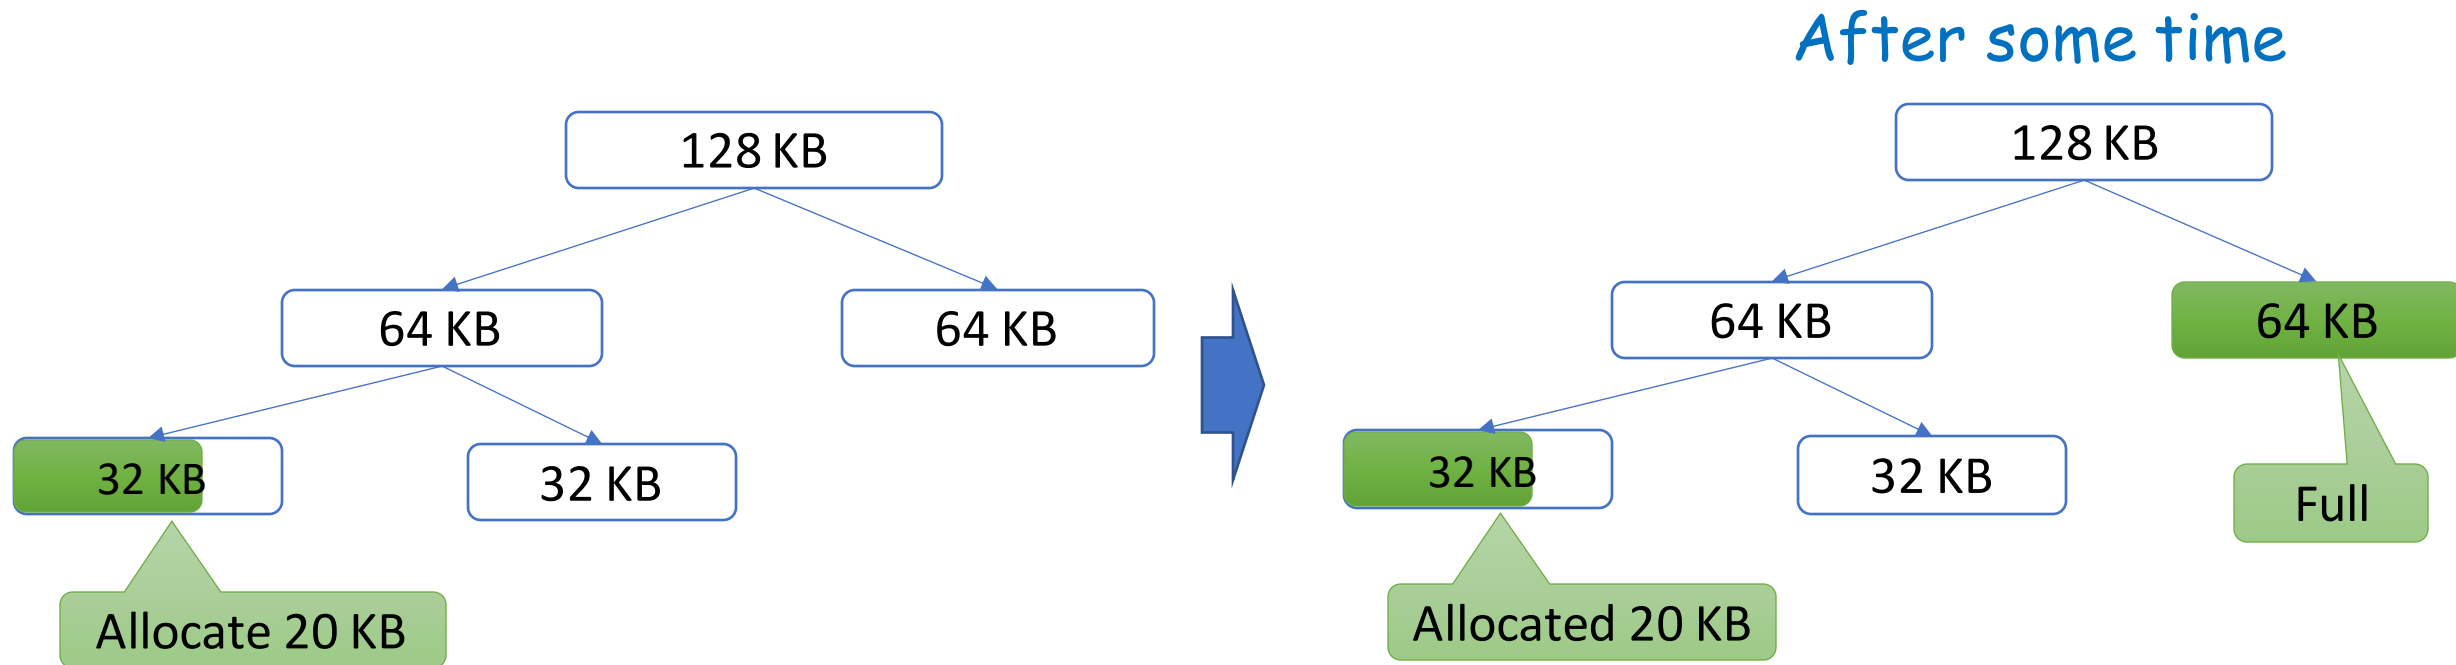
128 KB

Total memory size



How do we **allocate** a 20 KB chunk?


Allocation using the Buddy Allocator



Now **delete** the 20 KB **chunk**



Contd



- The **size** of each **buddy** chunk is a power of two
- Adjacent **free** chunks (holes) can be merged
- Tries to **minimize** internal and external **fragmentation**
- Can **handle** requests for both small and large **regions**




Create a **data structure** to effectively manage the **buddy list**

Implementation of the Buddy System



```
struct zone {  
    ...  
    struct free_area    free_area[MAX_ORDER];  
    ...  
}  
  
struct free_area {  
    struct list_head    free_list[MIGRATE_TYPES]; /* unmovable, movable,  
reclaimable, ...*/  
    unsigned long        nr_free;  
};
```



Traverse the *free_area* with an increasing order

 `/mm/page_alloc.c`

`__rmqueue_smallest(...)`



```
for (current_order = order; current_order < MAX_ORDER; ++current_order) {  
    area = &(zone->free_area[current_order]);  
    page = get_page_from_free_area(area, migratetype);  
    if (!page)  
        continue;  
    del_page_from_free_list(page, zone, current_order);  
    return page;  
}
```

- **Traverse** the list of **free** pages: **order**=0 to **order** = **MAX_ORDER** – 1
- Pick the first **free** compound page (from the first list that is non-empty)
- Delete the **selected page** from the corresponding **free** list



Merging *free* Pages

```
void __free_one_page(struct page *page, unsigned long pfn,
                    struct zone *zone, unsigned int order, ...)
{
    while (order < MAX_ORDER - 1) {
        buddy = find_buddy_page_pfn(page, pfn, order, &buddy_pfn);
        del_page_from_free_list(buddy, zone, order);

        ....
        combined_pfn = buddy_pfn & pfn;
        page = page + (combined_pfn - pfn);
        pfn = combined_pfn;
        order++;
    }

    set_buddy_order(page, order);
    add_to_free_list(page, zone, order, migratetype);
}
```

$buddy_pfn = pfn \wedge (1 \ll order)$

Not free any more

pfn of the parent

Details of $combined_pfn$

Set the order of the large combined_page and free it

Slab Allocator

- The **buddy** system is for generic memory allocation
- Most of the **time** objects of specific types are used. They have a **fixed** structure.
 - No **reason** to allocate and deallocate them.
 - Take up a large **memory** region and use it for **dedicated** object storage (of only one type)
 - For fast access, keep a **cache** of pre-initialized objects
 - **Use** and **return** → Similar to a **pool** of objects

Terms:

slab

cache

chunk

The Kernel has Three Object Allocators



SLOB Allocator

- Old Solaris-based allocator that uses the first-fit method.



SLAB Allocator

- Most popular allocator.

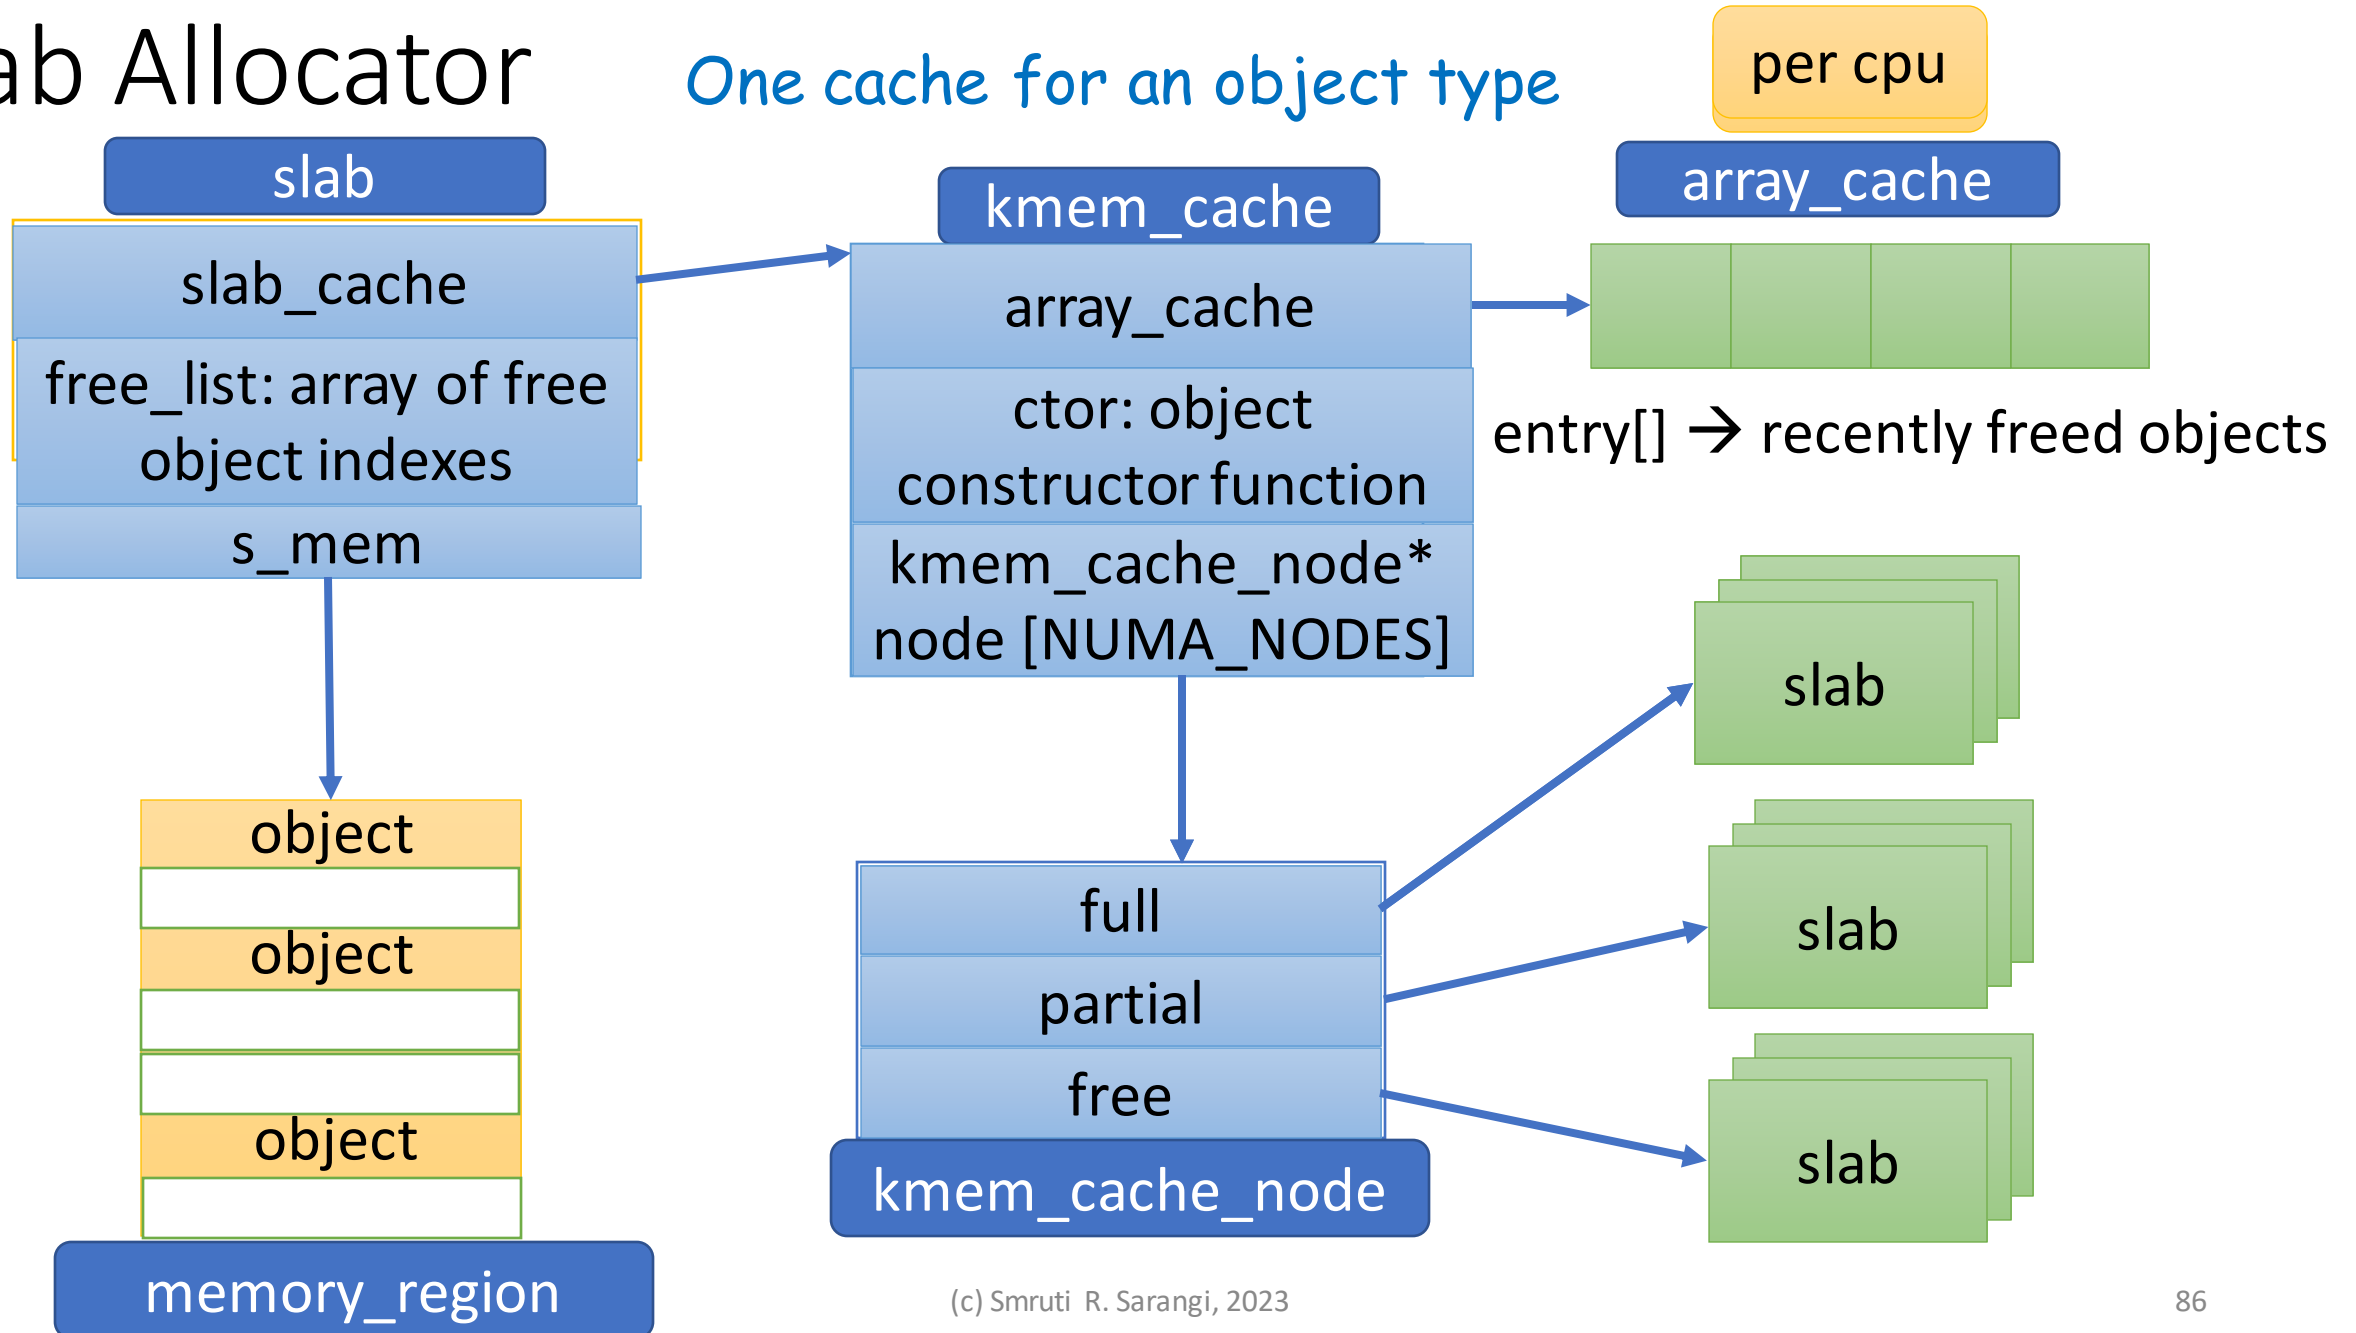


SLUB Allocator

- Better performance, scalability and simpler

Slab Allocator

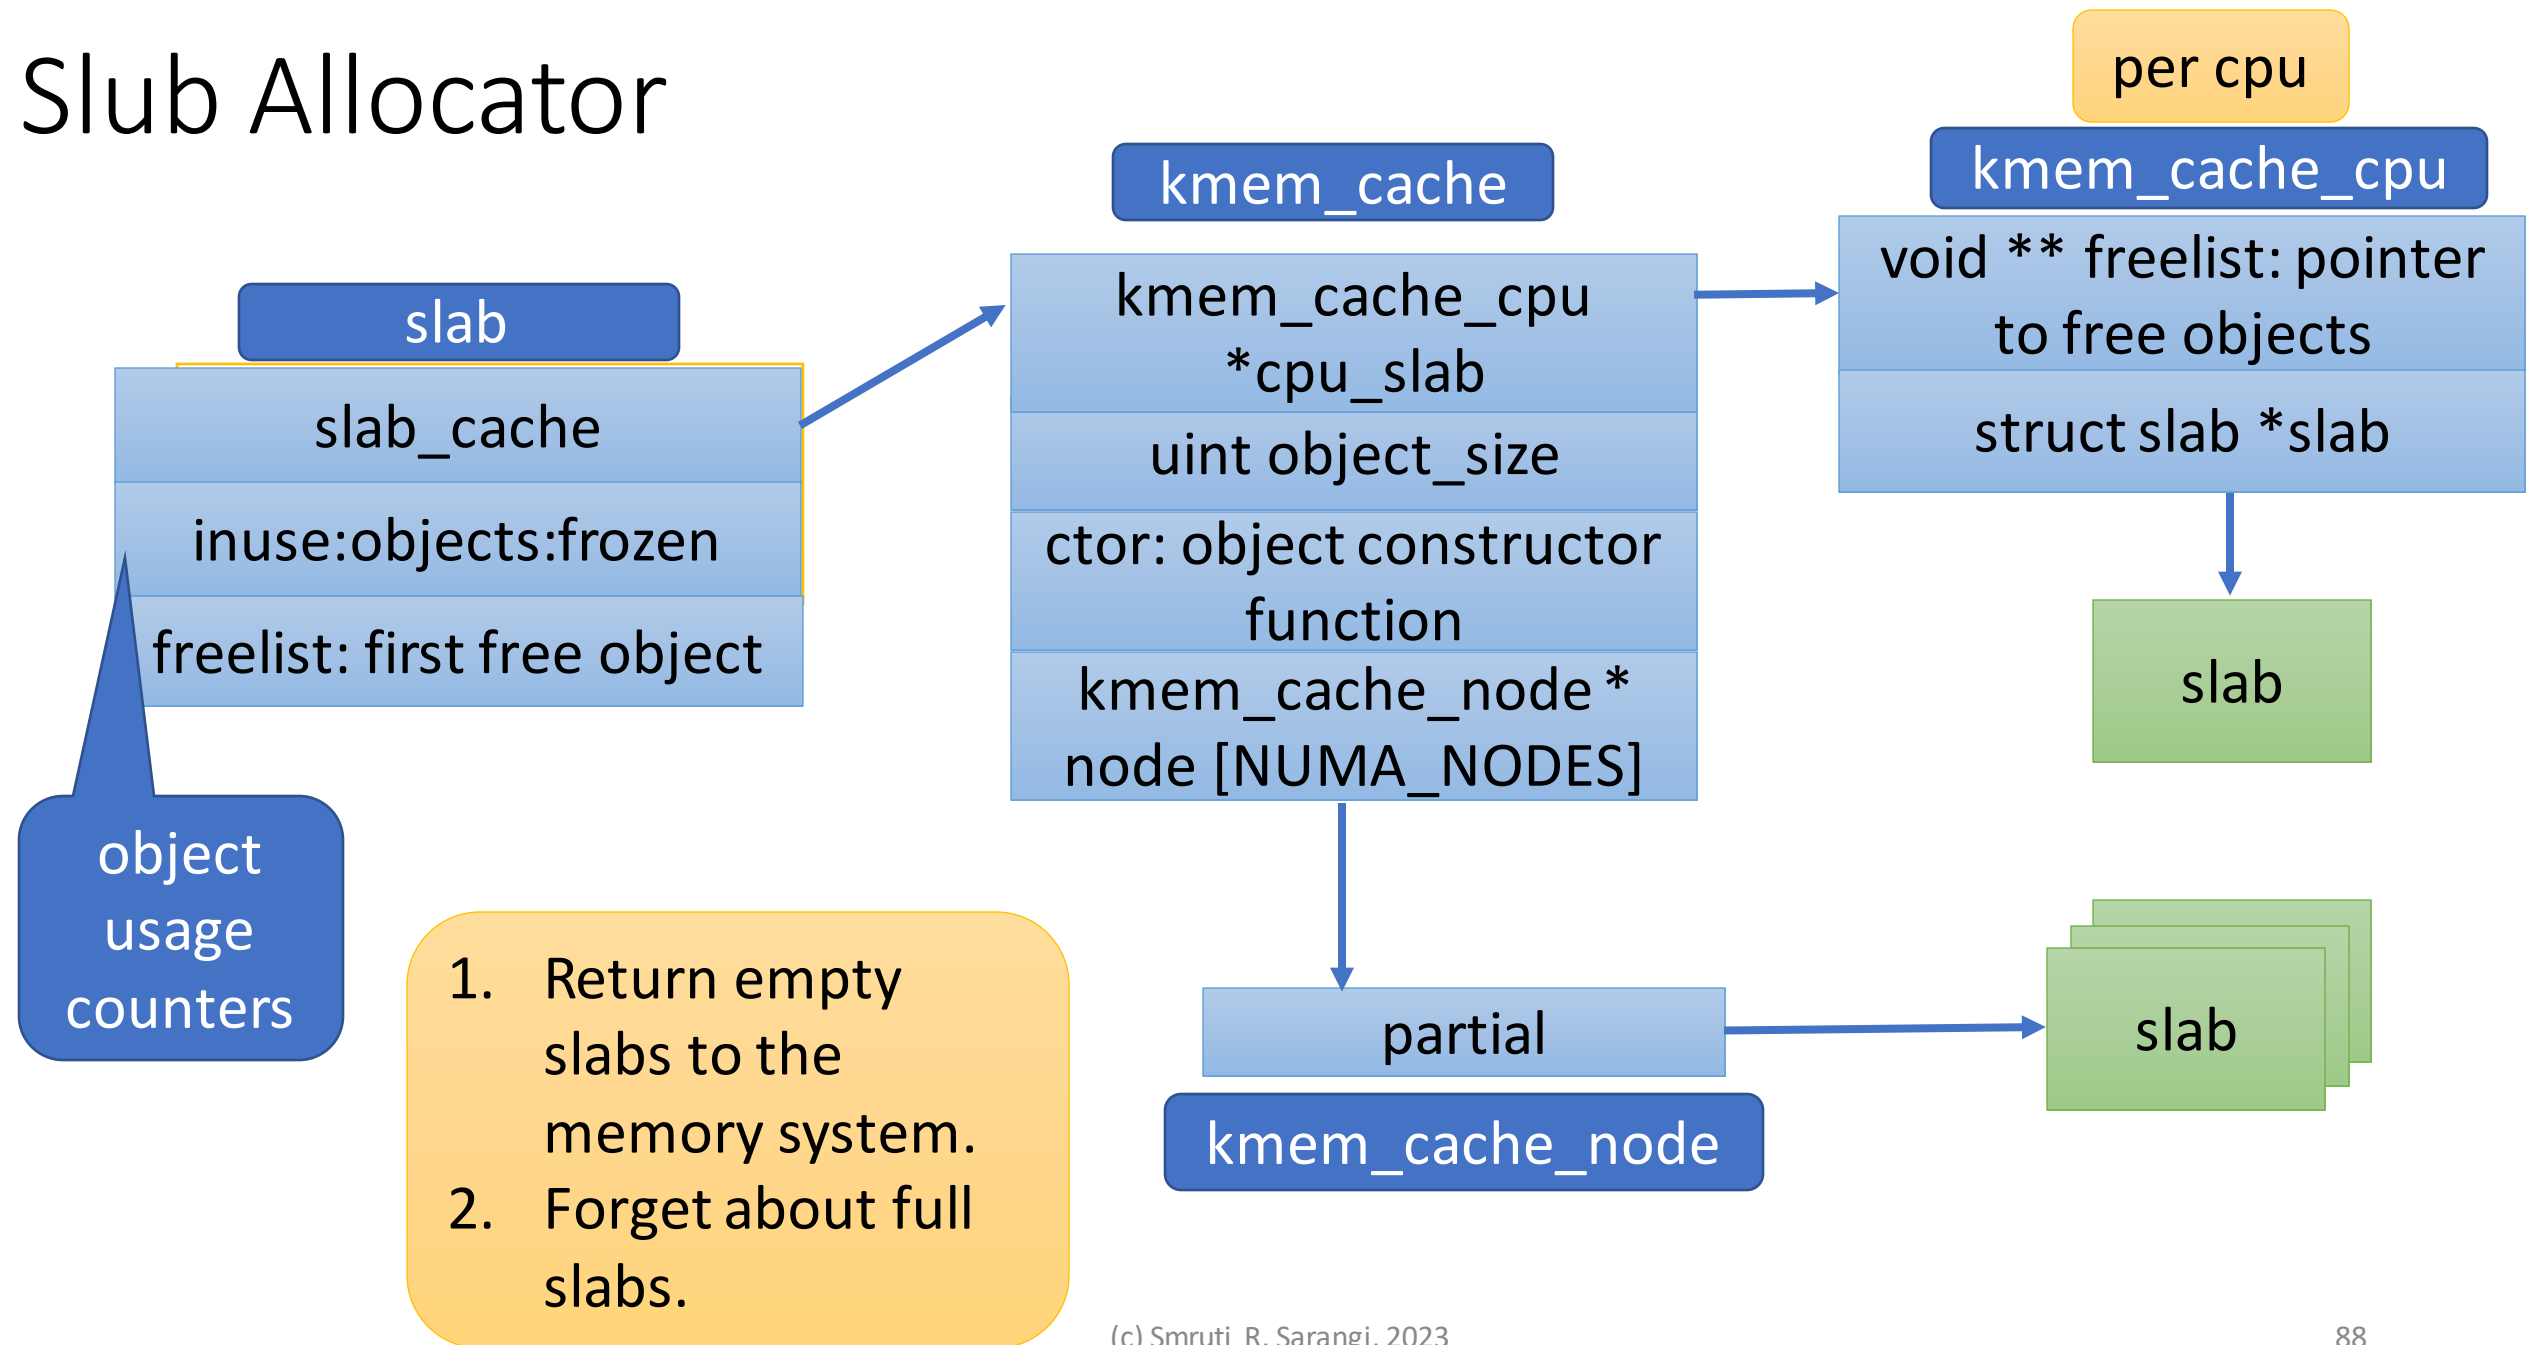
One cache for an object type



Quick Explanation of the SLAB scheme

- A **dedicated** region of the physical memory and the kernel's virtual memory stores slabs
- Every **slab** corresponds to a **region** of **physical** memory (set of **contiguous** physical pages)
- The *kmem_cache* is the **key** data structure **that** maintains all the slabs for a given **object type**. The external world only **interacts** with the *kmem_cache*
 - It maintains a list of **slabs** in three different states: **full**, **free**, and **partial**
 - This list is specific to each **NUMA** node
 - It also holds a **per-cpu cache** of **free** objects that can quickly be allocated
 - There are **few** *kmem_cache* **structures** in the system for storing very **frequently** used objects that are not allocated and deallocated every time they are used

Slab Allocator



Quick Explanation of the SLUB Allocator

- Much simpler than the SLAB allocator (still uses the notion of slabs)
- Here also, the *kmem_cache* is the key data structure that is visible to the external world
- In a slab → Maintain a freelist (object pointers)
- In the *kmem_cache* → Maintain an array of per-cpu slabs
 - One slab per CPU
 - Just maintain the list of partial slabs in a *kmem_cache* (a single list)
 - Basically means: forget about full and totally free slabs
 - free slabs are returned to the regular memory allocation system
 - When an object is freed in a full slab, add the full slab to the partial slab list

Bibliography

1. Belady, Laszlo A., Robert A. Nelson, and Gerald S. Shedler. "An anomaly in space-time characteristics of certain programs running in a paging machine." *Communications of the ACM* 12.6 (1969): 349-353.
2. FIFO anomaly is unbounded, [arXiv:1003.1336](https://arxiv.org/abs/1003.1336)



srsarangi@cse.iitd.ac.in

thank you

